

Intro to shaders and OpenGL



Realtime rendering intro

Vocabulary

Real-time rendering is the technology that allows to generate images so quickly that the human see them react immediatly to its inputs. Real-time rendering is rotoscopy, like in cinema, but with a quicker rate and without the need to pregenerate images.

Nowadays, most real-time rendering relies on the use of graphics cards (GPU), that contains hundreds if not thousands of small processors, specialized in 3d computing. Multiples computations are triggered each generated frames. Some of those computations are fixed and are sometimes implemented in the very hardware of the GPU. Others can be modified, either by giving them some parameters, or by providing complete programs.

Those custom programs are known as "Shaders". They are written in a specific "shader language", which looks like C but possess special types and functions that are handy for GPU instructions.

OpenGL is a standard API to address instructions to the GPU. It comes with a shader language called GLSL. OpenGL can be used on most desktop and living-room machines (Windows, Linux, consoles...) and, in its OpenGL ES version, on android phones. There is also a WebGL API, for browsers, that have some similarities with OpenGL. MacOS and iOS no longer support OpenGL and replace it with a specific API called Metal. Microsoft machines (PC Windows, XBox console) can also run a specific API called DirectX.

From 2016, a new API called Vulkan is meant to replace OpenGL. However, it is harder to setup and to understand, so we will only see it in 4th year.

Content of the lecture

We will take advantage of this lecture to :

- Learn compilation from scratch, without the help of the huge Visual Studio
- Discover the very basics of OpenGL and Shaders in 2D
- Build very small realtime engine in C++, with a scene system, logging and shader handling capabilities

You now understand this will be a quite technical course. You will feel a glimpse of the real taste of low level programming. This course is meant for people who want to become programmers. Nevertheless, it is accessible to other students, and will learn you a lot in term of general real-time or game culture.

Installation

MinGW as a compile tool

Under Linux, programmers use a compiler called g++. For this lecture, we don't want to use the standard windows compiler (vc), but the g++ compiler. We will use the g++ compiler, directly giving instructions to it.

If you are under Linux or MacOS, you just have to install the g++ package and ignore this page.

Install MinGW x64

Go to : <https://sourceforge.net/projects/mingw-w64/>

Download the file. Install it.

In the install options, change Architecture from i686 to x86_64. Then validate the rest of the install.

Your program shoud be installed in the C:\Program Files\mingw-w64\x86_64-8.1.0-posix-seh-rt_v6-rev0 folder. (Maybe the version number is different.)

Add it to your PATH environment variable

Once MinGW is installed, we want the tools it provides to be added to our windows command line.

To do that, there is a special windows environment variable called PATH, that allow to simulate a folder to be present in the current folder of the command line.

Click on the start menu and type "environment". You'll see a command "Modify system environment variables" appear. Click on it. On the tab use the last button "Environment variables..." on the bottom.

In the new window, you will see a line with "Path" at the beginning. Double click it. A new windows with multiple paths will appear. Click on new, then add the absolute path to the ./mingw-w64/bin folder that contains the MinGW executables. For instance, for my installation it would be :

C:\Program Files\mingw-w64\x86_64-8.1.0-posix-seh-rt_v6-rev0\mingw64\bin

Then click OK and close the windows. You can test if all went well by opening a new terminal and type "g++ --version". The command should work.

(Note that you must use a new terminal. If a previous terminal was open before you change the environment variable, the PATH won't be updated for this terminal.)

Handmade compilation

Hello handmade world

Folder structure

You need a folder for your project. Open it with visual code. Install the C++ plugin if needed.

Inside this folder, create a "src" folder and a "external" folder. The src will contain the source code and the external folder libraries we will use to build our application.

First program

Create a Main.cpp file in the src folder, with this content :

```
#include <iostream>

using std::cout;
using std::endl;

int main(int argc, char** argv) {
    cout << "Hello handmade world" << endl;

    return 0;
}
```

Compile and run

Open the visual code terminal, and run the command :

```
g++ src/main.cpp -o hello.exe
```

It should compile a hello.exe application. Run it with the terminal :

```
./hello.exe
```

Congratulations, you have compiled and run your first handmade program.

Makefile

Purpose

We will now use a makefile to configure our compilation. Writing your compile command in the terminal will prove difficult and redundant, and makefiles are a good way to automate part of the process.

Makefile

Create a script folder at the root of your project. Create a "makefile" file (no extension).

```
all: hello.exe

hello.exe: Main.o
    g++ -o hello.exe Main.o

Main.o: ../src/Main.cpp
    g++ -c ../src/Main.cpp -std=c++17 -g -Wall -Wextra -pedantic
```

Then go to the script folder with your terminal and run:

```
mingw32-make
```

If you are under linux, you can simply run "make".

This will compile your program following the rules of the makefile document. The rules are the following :

- When we launch make without indication, make the "all" target
- This all target means building the hello.exe target
- The hello.exe target has a dependency : the Main.o target
- If this dependency is found, the hello.exe target will run "g++ -o hello.exe Main.o" in the terminal
- the main.o target has a dependency: "../src/Main.cpp"
- If this dependency is found, run "g++ -c main.o -o main.o -std=c++17 -g -Wall -Wextra -pedantic" in the terminal

This will generate a main.o and a hello.exe file in the scripts folder. The Main.o file is known as an object file, and the exe file is an application.

g++ options

If we compare with the previous slide, we use new g++ options.

-o is used to specify an "object file", which is actually the compiled cpp file (properly named "translation unit")

-std=c++17 selects the version on c++ standard (here c++17) we use in our program

-g mean we want debug symbols, to allow debugging. The compilation is not optimized.

-Wall is a warning configuration command. We want all usual warnings.

-Wextra means we want additional warnings

-pedantic will check the code is conform to the c++ norm

More compilation options (french) : <https://openclassrooms.com/fr/courses/1262861-c-maitriser-le-compilateur-g>

Tidying our build files

We don't want to store our files in the scripts folder. Instead, we will create a build folder at the root, and a obj folder in the build folder. We can then modify our makefile :

```
all: hello.exe

hello.exe: Main.o
    g++ -o ./build/hello.exe ./build/obj/Main.o

main.o: ../src/Main.cpp
    g++ -c ../src/Main.cpp -std=c++17 -g -Wall -Wextra -pedantic
```

Now our exe go directly in the build folder and our obj files go into build/obj.

An automatic Makefile

Declare folders

We give by end the path to our translation units and object files. We In order to avoid mistakes, we can store those paths in variables

makefile

```
PROJECT := hello.exe
SRC_DIR := ./src
BUILD_DIR := ./build
OBJ_DIR := ./build/obj
EXT_DIR := ./external

all: $(PROJECT)

$(PROJECT): main.o
    g++ -o $(BUILD_DIR)/$(PROJECT) $(OBJ_DIR)/main.o

main.o: $(SRC_DIR)/main.cpp
    g++ -o $(OBJ_DIR)/main.o -c $(SRC_DIR)/main.cpp -std=c++17 -g -Wall -Wextra -pedantic
```

Compiling multiple folders

Our code will be divided into multiple folders. In this project, we will have an engine folder and a game folder into the src folder. Create them.

Also add a Window.h, Game.h and Window.cpp file in the engine folder, and a Game.cpp file in the game folder, with this content :

Window.h

```
#ifndef WINDOW_H
#define WINDOW_H

class Window {
public:
    Window();
    void display();
};

#endif
```

Window.cpp

```
#include "window.h"
#include <iostream>

using std::cout;
using std::endl;

Window::Window() {}

void Window::display() {
    cout << "Window is displaying" << endl;
}
```

Game.h

```
#ifndef GAME_H
#define GAME_H

class Game {
public:
    Game();
    void display();
};

#endif
```

Game.cpp

```
#include "../engine/Game.h"
#include <iostream>

using std::cout;
using std::endl;

Game::Game() {}

void Game::display() {
    cout << "Game is displaying" << endl;
}
```

Update Main.cpp to use our new classes :

```
#include <iostream>

#include "engine/Game.h"
#include "engine/Window.h"

using std::cout;
using std::endl;

int main(int argc = 0, char **argv = nullptr) {
    if (argc > 0) {
        for (int i = 0; i < argc; ++i) {
            cout << argv[i] << endl;
        }
    }

    Window window;
    Game game;

    cout << "Hello handmade world" << endl;
    window.display();
    game.display();

    return 0;
}
```

Now you will need to add a game and engine folder in the build/obj folder, then to update the makefile in order to auto-detect cpp files :

```
PROJECT := hello.exe
SRC_DIR := ./src
BUILD_DIR := ./build
OBJ_DIR := ./build/obj
EXT_DIR := ./external

SRC_FILES := $(wildcard $(SRC_DIR)/**/*.*)
OBJ_FILES := $(patsubst $(SRC_DIR)/%.*, $(OBJ_DIR)/%.o, $(SRC_FILES))

all: $(PROJECT)

$(PROJECT): $(OBJ_FILES)
    g++ -g -o $(BUILD_DIR)/$@ $^

$(OBJ_DIR)/%.o: $(SRC_DIR)/%.*
    g++ -std=c++17 -g -Wall -Wextra -pedantic -c -o $@ $<
```

Scripts to improve usability (1/3)

The assets script

Imagine we have data (images, levels...) in our game. How would we handle them? How would the compiled program retrieve them?

First, let's create an "assets" folder in the root of the project. This will be the folder that will contain our data. Add a text.txt file inside, so we can test it.

When we run the program, we will suppose the assets folder will resides in the same folder as the exe file. Thus, we need to copy it in the build folder. We will use a script to do that.

Create a assets.bat file (linux : assets.sh) in the scripts folder :

```
set buildDir=%~dp0..\build
set assetsDir=%~dp0..\assets
set extDir=%~dp0..\external

:: Copy assets
if not exist %buildDir%\assets mkdir %buildDir%\assets
xcopy /y /s %assetsDir% %buildDir%\assets
```

Under linux, assets.sh :

```
#!/bin/bash

dot=$(dirname "$0")
buildDir=$dot../build
assetsDir=$dot../assets

# Copy assets
if [ ! -d "$buildDir/assets" ]; then
    mkdir "$buildDir/assets"
fi
cp -a "$assetsDir/" "$buildDir/assets"
```

You can test the script the same way you run a program in the terminal.

Scripts to improve usability (2/3)

The clean script

For now, we have to manually delete our files if we want to start over the build. Let's create a clean.bat script (linux : clean.sh) in the scripts folder.

clean.bat

```
@echo off

set buildDir=%~dp0..\build
set objDir=.\obj\
set assetDir=.\assets\

if exist %buildDir% (
    pushd %buildDir%
    del /q /s *.exe *.pdb *.ilk *.dll
    rd /s /q %objDir%
    if exist %assetDir% rd /s /q %assetDir%
    popd
)
```

Under linux, clean.sh :

```
#!/bin/bash

dot="$(pwd)/$(dirname "$0")"
buildDir=$dot/../build
buildObjDir=$buildDir/obj/
assetDir=$dot/../assets/

if [ -d "$buildDir" ]; then
    cd $buildDir
    rm *.exe *.pdb *.ilk *.dll
    rm -r $buildObjDir
    if [ -d "$buildDir/assets" ]; then
        rm -r "$buildDir/assets"
    fi
fi
```

Running this script will delete the exe file and remove all build/obj folders.

The problem is we cannot compile anymore : we need the obj folder and subfolder for our makefile to do its job.

Scripts to improve usability (3/3)

The build script

Our build script will create the folders we need in the build folder and launch our makefile automatically.

Create a build.bat file :

```
@echo off

:: Create build dir
set buildDir=%~dp0..\build
if not exist %buildDir% mkdir %buildDir%
pushd %buildDir%

:: Create obj dir
set objDir=.\obj
if not exist %objDir% (
    mkdir %objDir%
    mkdir %objDir%\engine
    mkdir %objDir%\game
)

:: Needed folder
set scriptDir=%~dp0..\scripts

:: Use make to build default target
cd %scriptDir%\mingw32-make

popd
```

Under linux, build.sh :

```
#!/bin/bash

# Create build dir
dot="$(pwd)=$(dirname "$0")"
buildDir=$dot/../build
if [ ! -d "$buildDir" ]; then
    mkdir $buildDir
fi

# Create obj dir
objDir="$buildDir/obj"
if [ ! -d "$objDir" ]; then
    mkdir $objDir
    mkdir $objDir/engine
    mkdir $objDir/game
fi

# Needed folder
scriptDir=$dot/..../scripts

# Use make to build default target
cd $scriptDir
make

cd $dot
```

Shortcuts to launch scripts (1/2)

Build with Ctrl + Shift + B

By default, visual code provides a Ctrl + Shift + B shortcut to build. We want a specific action on this shortcut : launching the build script.

Use Ctrl + Shift + P to open the command panel, and choose Config default build task.

Choose any options and replace the content of ./vscode/tasks.json with (works for windows and unix) :

```
{
  "version": "2.0.0",
  "tasks": [
    {
      "label": "build",
      "type": "shell",
      "windows": {
        "command": "./scripts/build",
        "args": [
          ""
        ]
      },
      "linux": {
        "command": "./scripts/build.sh"
      },
      "group": {
        "kind": "build",
        "isDefault": true
      },
      "problemMatcher": {
        "owner": "cpp",
        "fileLocation": [
          "relative",
          "${workspaceRoot}"
        ],
        "pattern": {
          "regexp": "^(.*):\\d+:\\d+/s+(warning|error):/s+(.*)$",
          "file": 1,
          "line": 2,
          "column": 3,
          "severity": 4,
          "message": 5
        }
      }
    },
    {
      "label": "clean",
      "type": "shell",
      "windows": {
        "command": "./scripts/clean",
        "args": [
          ""
        ]
      },
      "linux": {
        "command": "./scripts/clean.sh"
      }
    },
    {
      "label": "prepare-assets",
      "type": "shell",
      "windows": {
        "command": "./scripts/assets",
        "args": [
          ""
        ]
      },
      "linux": {
        "command": "./scripts/assets.sh"
      }
    }
  ]
}
```

Shortcuts to launch scripts (2/2)

Clean with Ctrl + Shift + C

There is no default shortcut to clean, so we will create a new one. Use Ctrl + Shift + P then Open Keyboard shortcut (json).

Put this in your personal bindings :

```
[  
  {  
    "key": "ctrl+shift+c",  
    "command": "workbench.action.tasks.runTask",  
    "args": "clean"  
  }  
]
```

The clean task is already set up thanks to the previous slide.

Launch debug

The launch configuration

By default in visual code, the F5 key is associated to debug launch. Because we don't use the usual windows build, we cannot auto-create our debug launch configuration.

Create a ./vscode/launch.json file and fill it with (also works for linux) :

```
{
  "version": "0.2.0",
  "configurations": [
    {
      "name": "(Windows g++) Launch",
      "type": "cppdbg",
      "request": "launch",
      // This is the name of your executable
      "program": "${workspaceRoot}/build/hello.exe",
      "args": [],
      "stopAtEntry": false,
      "cwd": "${workspaceFolder}/build",
      "environment": [],
      "externalConsole": true,
      "MIMode": "gdb",
      // This is your path to mingw's gdb
      "miDebuggerPath": "C:/Program Files/mingw-w64/x86_64-8.1.0-posix-seh-rt_v6-rev0/mingw64/bin/gdb.exe",
      "setupCommands": [
        {
          "description": "Enable pretty-printing for gdb",
          "text": "-enable-pretty-printing",
          "ignoreFailures": true
        }
      ],
      // This will execute the assets script before launch
      "preLaunchTask": "prepare-assets"
    },
    {
      "name": "(Linux g++) Launch",
      "type": "cppdbg",
      "request": "launch",
      "program": "${workspaceRoot}/build/hello",
      "args": [],
      "stopAtEntry": false,
      "cwd": "${workspaceFolder}/build",
      "environment": [],
      "externalConsole": true,
      "MIMode": "gdb",
      "setupCommands": [
        {
          "description": "Enable pretty-printing for gdb",
          "text": "-enable-pretty-printing",
          "ignoreFailures": true
        }
      ],
      "preLaunchTask": "prepare-assets"
    }
  ]
}
```

You can check the assets script was launched.

MinGW uses a debugger called gdb, which is the debugger associated with g++. Here we launch our debug build with this program.

Test the debugger

Add a breakpoint into your main, by left clicking on the left of the line number, and run the program with F5. The program should stop right at your breakpoint. Try the different debug features.

Exercice : the release build

Exercice

For now, we only build and launch a debug build. A debug build is not optimized, and slower than a release build.

In this exercice, you will create a release build for our small program and update all the others scripts and config files. Here is some information:

- Release build will be built into a ./release folder
- Shortcut to run the release build will be Ctrl + Shift + R
- Release build needs assets and shall be cleaned by the clean script

You need to know the g++ command to compile with optimisations. For the release target :

```
g++ -O2 -mwindows -o $(RELEASE_DIR)/$(PROJECT) $^
```

For each translation units :

```
g++ -std=c++17 -O2 -Wall -Wextra -pedantic -c -o $@ $<
```

The -O2 parameter give the compiler the optimization level. -O3 is actually the most optimized build, but it happens that -O2 is faster.

Going further

A good resource about compilation optimization:

<https://www.lihaoyi.com/post/HowanOptimizingCompilerWorks.html>

Setting up SDL2

Downloading SDL2

SDL is a library that allows to handle most common game operations : running a window, handling inputs, setting up the display. It possesses extensions to add image format, sound, text...

Download SDL2 Development libraries from <https://www.libsdl.org/download-2.0.php>

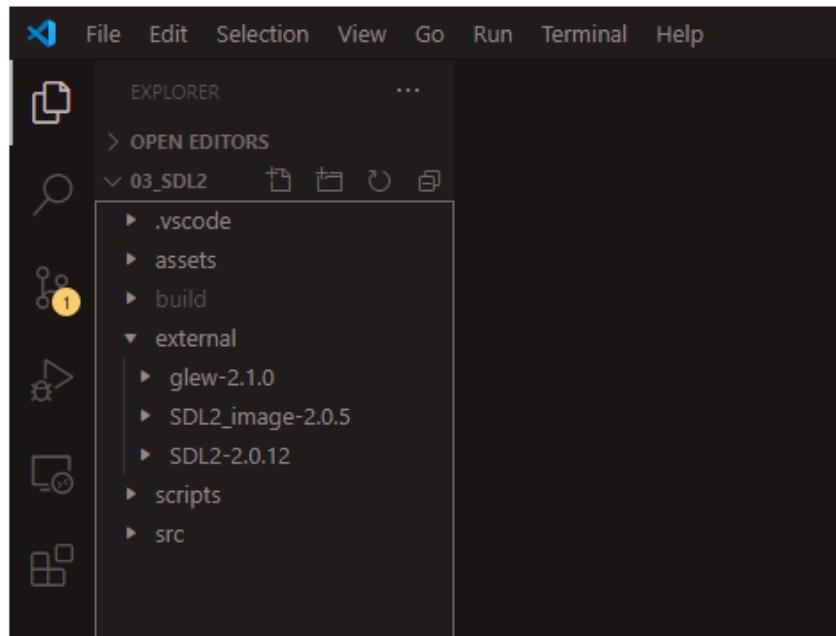
You have to choose `SDL2-devel-2.0.x-VC.zip` (Visual C++ 32/64-bit) - just after the middle of the page. Unzip it to your external folder.

Do the same with `SDL_image` : https://www.libsdl.org/projects/SDL_image/. You have to choose `SDL2_image-devel-2.0.5-VC.zip` (Visual C++ 32/64-bit). Also unzip it to the external folder.

Finally download `glew` : <https://sourceforge.net/projects/glew/files/glew/2.1.0/glew-2.1.0-win32.zip/download>

Unzip it to the same folder.

You should have a folder structure like this one :



Under linux, choose the right version.

Configuring includes

We need our code and our intellisense to find SDL and GLEW files, so that we can use them in code.

To allow VS Code to find our libraries, hit Ctrl + Shift + P then C++ Edit configuration (JSON). Change the code so that it looks like this :

```
{
  "configurations": [
    {
      "name": "Win32",
      "includePath": [
        "${workspaceFolder}/**",
        "${workspaceRoot}/external/SDL2-2.0.12/include",
        "${workspaceRoot}/external/SDL2_image-2.0.5/include",
        "${workspaceRoot}/external/glew-2.1.0/include"
      ],
      "defines": [
        "_DEBUG",
        "UNICODE",
        "_UNICODE"
      ],
      "windowsSdkVersion": "10.0.18362.0",
      "compilerPath": "C:/Program Files (x86)/Microsoft Visual Studio/2019/Community/VC/Tools/MSVC/14.28.29333/bin/Hostx64/x64/cl.exe",
      "cStandard": "c17",
      "cppStandard": "c++17",
      "intelliSenseMode": "msvc-x64"
    }
  ],
  "version": 4
}
```

Now, on the top of your main.cpp file, you should add without error :

```
#include <SDL.h>
```

```
...
```

Update the makefile

Using libraries implies compiling with them. We need to update the makefile to tell it to use libraries.

Fortunately, variables come to our help once more :

```
PROJECT := hello.exe
SRC_DIR := ./src
BUILD_DIR := ./build
OBJ_DIR := ./build/obj
EXT_DIR := ./external

SRC_FILES := $(wildcard $(SRC_DIR)/**/*.cpp) $(wildcard $(SRC_DIR)/*/*.cpp)
OBJ_FILES := $(patsubst $(SRC_DIR)/%.cpp,$(OBJ_DIR)/%.o,$(SRC_FILES))

LIBRARIES := -ISDL2main -ISDL2 -ISDL2_image -Iglew32 -Izlib1 -lopengl32

INCLUDE :=-I$(EXT_DIR)\SDL2-2.0.12\include \
           -I$(EXT_DIR)\SDL2_image-2.0.5\include \
           -I$(EXT_DIR)\glew-2.1.0\include

LIB :=-L$(EXT_DIR)\SDL2-2.0.12\lib\x64 \
      -L$(EXT_DIR)\SDL2_image-2.0.5\lib\x64 \
      -L$(EXT_DIR)\glew-2.1.0\lib\Release\x64

all: $(PROJECT)

$(PROJECT): $(OBJ_FILES)
    g++ -g -o $(BUILD_DIR)/$@ $^ $(LIB) $(LIBRARIES)

$(OBJ_DIR)/%.o: $(SRC_DIR)/%.cpp
    g++ -std=c++17 -g -Wall -Wextra -pedantic -c -o $@ $< $(INCLUDE)
```

For linux, you just have to update :

```
LIBRARIES := -ISDL2main -ISDL2 -ISDL2_image -IGLEW -IGLU -IGL

$(PROJECT): $(OBJ_FILES)
    g++ -g -o $(BUILD_DIR)/$@ $^ $(LIBRARIES)
```

Update assets script

Finally, our program need a compiled version of the librairies to run. Under windows, those compiled version are known as dll files.

We shall copy the librairies dll to the build folder when we will copy assets.

Let's update the assets.bat script :

```
set buildDir=%~dp0..\build
set assetsDir=%~dp0..\assets
set extDir=%~dp0..\external

:: Copy dependencies
if not exist %buildDir%\SDL2.dll xcopy /y %extDir%\SDL2-2.0.12\lib\x64\SDL2.dll %buildDir%
if not exist %buildDir%\SDL2_image.dll xcopy /y %extDir%\SDL2_image-2.0.5\lib\x64\SDL2_image.dll %buildDir%
if not exist %buildDir%\libjpeg-9.dll xcopy /y %extDir%\SDL2_image-2.0.5\lib\x64\libjpeg-9.dll %buildDir%
if not exist %buildDir%\libpng16-16.dll xcopy /y %extDir%\SDL2_image-2.0.5\lib\x64\libpng16-16.dll %buildDir%
if not exist %buildDir%\zlib1.dll xcopy /y %extDir%\SDL2_image-2.0.5\lib\x64\zlib1.dll %buildDir%
if not exist %buildDir%\glew32.dll xcopy /y %extDir%\glew-2.1.0\bin\Release\x64\glew32.dll %buildDir%

:: Copy assets
if not exist %buildDir%\assets mkdir %buildDir%\assets
xcopy /y /s %assetsDir% %buildDir%\assets
```

Under linux, you do not need to copy the files.

A first window

A first window

Boilerplate code to init everything :

```
#include <SDL.h>
#include <GL/glew.h>
#include <iostream>

using std::cout;
using std::endl;

constexpr int SCREEN_WIDTH = 640;
constexpr int SCREEN_HEIGHT = 480;

int main(int argc = 0, char **argv = nullptr) {
    // Handle args
    if (argc > 0) {
        for (int i = 0; i < argc; ++i) {
            cout << argv[i] << endl;
        }
    }

    SDL_Window* window = nullptr;
    SDL_GLContext context;
    int flags = SDL_WINDOW_OPENGL;

    if (SDL_Init(SDL_INIT_EVERYTHING) != 0) {
        cout << "SDL failed to initialize" << endl;
        return 1;
    }

    // Initialize window and openGL
    window = SDL_CreateWindow("First triangle", SDL_WINDOWPOS_CENTERED, SDL_WINDOWPOS_CENTERED, SCREEN_WIDTH, SCREEN_HEIGHT, flags);
    context = SDL_GL_CreateContext(window);
    glewExperimental = GL_TRUE;
    glewInit(); ← Start GLEW then config OpenGL

    // Get info
    const GLubyte* renderer = glGetString(GL_RENDERER);
    const GLubyte* version = glGetString(GL_VERSION);
    cout << "Renderer: " << renderer << endl;
    cout << "OpenGL version supported: " << version << endl;

    // Tell GL to only draw onto a pixel if the shape is closer to the viewer
    glEnable(GL_DEPTH_TEST); // enable depth-testing
    glDepthFunc(GL_LESS); // depth-testing interprets a smaller value as "closer"

    // Set viewport and clear color
    glViewport(0, 0, SCREEN_WIDTH, SCREEN_HEIGHT);
    glClearColor(0.0f, 0.0f, 0.2f, 1.0f);

    // Load ← Spot this, because we will code here
    // Game loop
    bool isRunning = true;
    while(isRunning) {
        // Inputs
        SDL_Event event;
        while (SDL_PollEvent(&event)) {
            switch (event.type) {
                case SDL_QUIT:
                    isRunning = false; ← Needed to quit when pressing the cross
                    break;

                default:
                    break;
            }
        }
        // Update

        // Draw
        glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT); // Clear the screen
        /* Draw Here */ ← We will also code here

        SDL_GL_SwapWindow(window); // Swapbuffer
    }

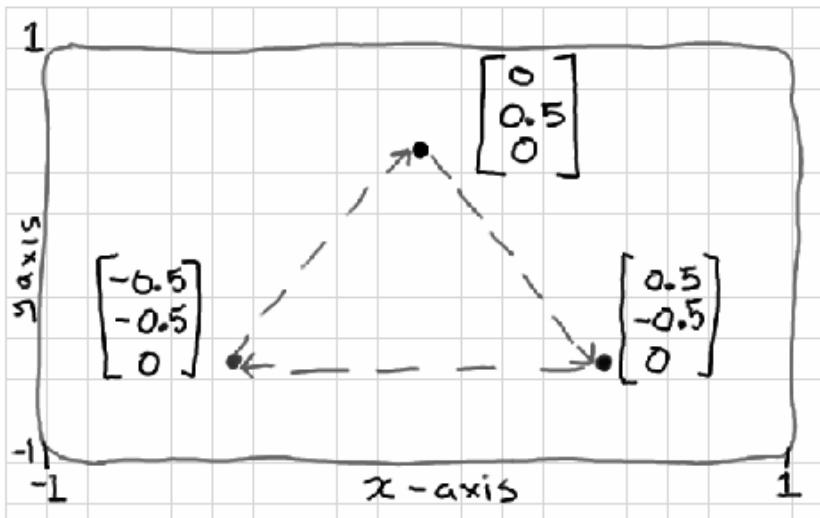
    // Quit
    SDL_DestroyWindow(window);
    SDL_GL_DeleteContext(context);

    return 0;
}
```

Display a triangle on your computer

Define a triangle in the vertex buffer

A triangle in screen view coordinate system



A triangle in a VBO (Vertex Buffer Object)

We pack all the triangle points, 3 by 3 in clockwise order, in a float array. Then we copy this array onto the graphics card using a VBO.

A VBO is a way to store data before sending it to the GPU.

```
...  
// Load  
GLfloat points[] = {  
    0.0f, 0.5f, 0.0f,  
    0.5f, -0.5f, 0.0f,  
    -0.5f, -0.5f, 0.0f  
};  
GLuint vbo = 0;  
glGenBuffers(1, &vbo);  
 glBindBuffer(GL_ARRAY_BUFFER, vbo);  
 glBufferData(GL_ARRAY_BUFFER, sizeof(points), points, GL_STATIC_DRAW);
```

Generate an empty buffer
Set it as the current buffer in OpenGL state machine by binding
Copy the points on the bound buffer

A mesh in the VAO (Vertex Attribute Object)

Most meshes will use a collection of one or more vertex buffer objects to hold vertex points, texture-coordinates, vertex normals, etc. In older GL implementations we would have to bind each one, and define their memory layout, every time that we draw the mesh.

To simplify that, we have a thing called the vertex attribute object (VAO), which remembers all of the vertex buffers that you want to use, and the memory layout of each one. We set up the vertex array object once per mesh. When we want to draw, all we do then is bind the VAO and draw.

```
...  
GLuint vao = 0;  
 glGenVertexArrays(1, &vao);  
 glBindVertexArray(vao);  
 glEnableVertexAttribArray(0);  
 glBindBuffer(GL_ARRAY_BUFFER, vbo);  
 glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 0, NULL);
```

Generate a new VAO
Bind in focus in the OpenGL state machine
We have one VBO, so the attribute location will be 0
Define the layout for attribute 0, the variable are made of 3

Setting up the shaders

The shaders will define how to draw our shape from the VAO. We need at least a vertex shader (shader the points' positions of the mesh) and a fragment shader (shader to compute the pixels). For now shaders are simple, so we write them in a text variable.

Vertex and fragment shader

We pack all the triangle points, 3 by 3 in clockwise order, in a float array. Then we copy this array onto the graphics card using a VBO.

```
...  
const char* vertexShader =  
"#version 410\n"  
"in vec3 vp;"  
"void main () {"  
" gl_Position = vec4 (vp, 1.0);"  
"}";  
  
const char* fragmentShader =  
"#version 410\n"  
"out vec4 frag_colour;"  
"void main () {"  
" frag_colour = vec4 (0.8, 0.5, 0.0, 1.0);"  
"}";
```

The diagram shows the shader code with several annotations:

- A double-headed arrow between the "#version 410\n" line and the text "OpenGL Version".
- A double-headed arrow between the "in vec3 vp;" line and the text "Input variable is a vec3, which match the position of a point from the VAO".
- A double-headed arrow between the "gl_Position = vec4 (vp, 1.0);" line and the text "We output gl_Position, which is the point with a 1 as a 4th coordinate (homogeneous w)".
- A double-headed arrow between the "out vec4 frag_colour;" line and the text "We output a color in RGBA format, for each of the pixels that belong to the shape".
- A double-headed arrow between the "frag_colour = vec4 (0.8, 0.5, 0.0, 1.0);" line and the text "The color will be a nice golden yellow".

Using the shaders

Before using the shaders we have to load the strings into a GL shader, and compile them. Then we put compiled shaders into a single program.

```
...  
GLuint vs = glCreateShader (GL_VERTEX_SHADER);  
glShaderSource (vs, 1, &vertexShader, NULL);  
glCompileShader (vs);  
GLuint fs = glCreateShader (GL_FRAGMENT_SHADER);  
glShaderSource (fs, 1, &fragmentShader, NULL);  
glCompileShader (fs);  
  
GLuint shaderProgram = glCreateProgram ();  
glAttachShader (shaderProgram, fs);  
glAttachShader (shaderProgram, vs);  
glLinkProgram (shaderProgram);
```

The diagram shows the OpenGL shader compilation and linking process with annotations:

- A double-headed arrow between the "glCreateShader(GL_VERTEX_SHADER)" line and the text "Compile vextex shader".
- A double-headed arrow between the "glCreateShader(GL_FRAGMENT_SHADER)" line and the text "Compile fragment shader".
- A double-headed arrow between the "glCreateProgram()" line and the text "Create empty program".
- A double-headed arrow between the "glLinkProgram(shaderProgram)" line and the text "Attach the shaders".
- A double-headed arrow between the "glLinkProgram(shaderProgram)" line and the text "Link program".

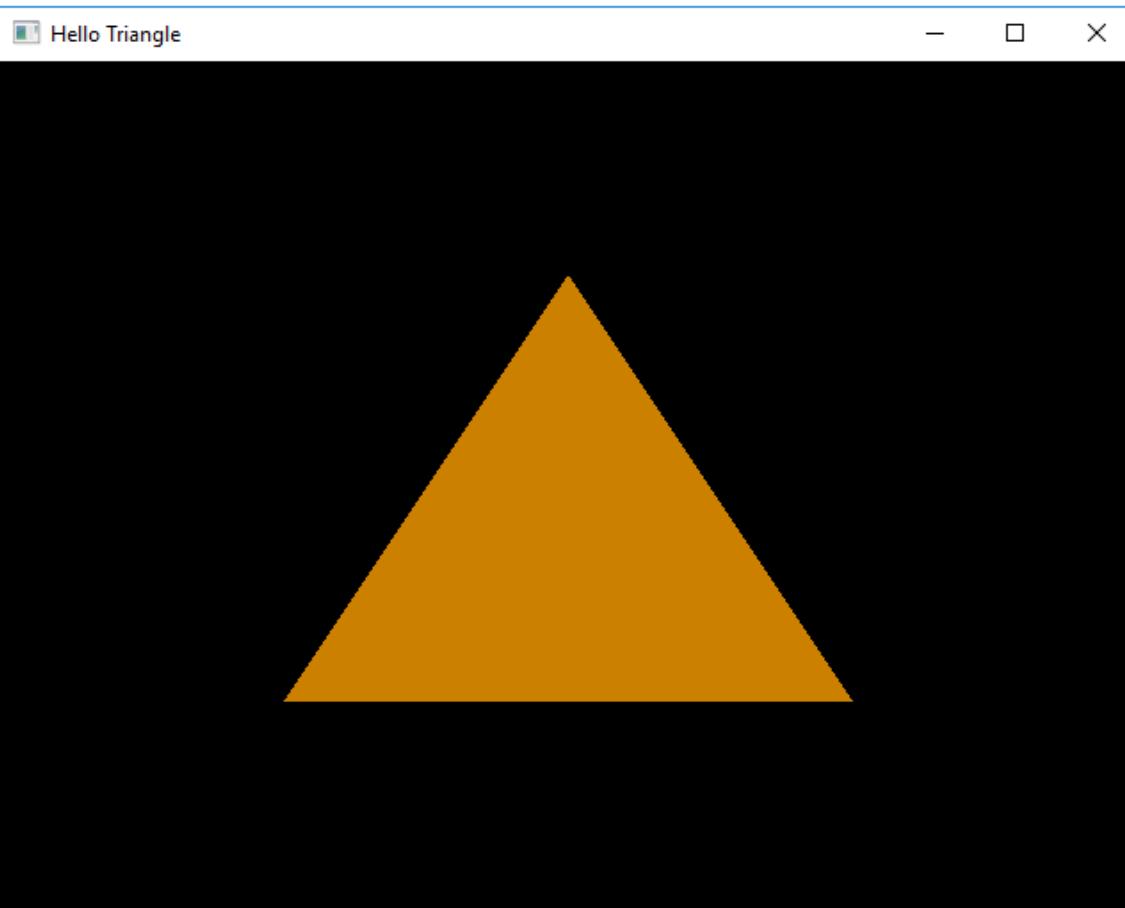
Drawing the triangle

In real time rendering, we clear the screen and redraw the scene each frame.

```
...  
    // Draw  
    glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT); // Clear the screen  
  
    /* Draw Here */  
    glUseProgram(shader_programme);  
    glBindVertexArray(vao);  
    glDrawArrays(GL_TRIANGLES, 0, 3);  
  
    SDL_GL_SwapWindow(window); // Swapbuffer  
...
```

Use program

draw points 0-3 from the currently bound VAO with current in-use shader



Exercises

- Load the shader strings from text files called test.vert and test.frag (a naming convention is handy).
- Change the colour of the triangle in the fragment shader.
- Try to move the shape in the vertex shader e.g. `vec4 (vp.x, vp.y + 1.0, vp.z, 1.0);`
- Try to add another triangle to the list of points and make a square shape. You will have to change several variables when setting up the buffer and drawing the shape. Which variables do you need to keep track of for each triangle? (hint: not much...).
- Try drawing with `GL_LINE_STRIP` or `GL_LINES` or `GL_POINTS` instead of triangles. Does it put the lines where you expect? How big are the points by default?
- Try changing the background colour by using `glClearColor ()` before the rendering loop. Something grey-ish is usually fairly neutral; `0.6f, 0.6f, 0.8f, 1.0f`.
- Try creating a second VAO, and drawing 2 shapes (remember to bind the second VAO before drawing again).
- Try creating a second shader programme, and draw the second shape a different colour (remember to "use" the second shader programme before drawing again).

Extended initialisation

A Logging system

Debuging graphical software is difficult, so we will start by implementing a Logging system.

Use cases

- Our logging system will log several log lines in a .log file
- It will have different level of logging
- It will have a config system that will allow us to choose which logging precision we want, and choose if we want to restart the logging
- Here is the logging use case :

```
LOG(Info) << "OpenGL version supported " << version;
```

- Here is the config setup :

```
LogConfig LOG_CONFIG = {};  
LOG_CONFIG.reporting_level = Debug;  
LOG_CONFIG.restart = true;
```

A Logging system

Header

src/engine/Log.h

```
// Inspired from http://www.drdobbs.com/cpp/logging-in-c/201804215
// and https://github.com/zuhd-org/easyloggingpp

#ifndef LOG_H
#define LOG_H

#include <time.h>
#include <stdio.h>
#include <fstream>
#include <sstream>

#define GL_LOG_FILE "gl.log" ← Log file name

enum LogLevel { Error, Warning, Info, Debug };

struct LogConfig {
    LogLevel reporting_level = Info;
    bool restart = false;
};

extern LogConfig LOG_CONFIG; ← We will use this global variable to store the config

class Log {
public:
    Log();
    virtual ~Log();
    std::ostringstream& get(LogLevel level = Info);
    static void restart(); ← restart will clear the log file, static function

private:
    std::ostringstream os;
    static std::ofstream file; ← restart need the file to be static too

    std::string get_label(LogLevel type);

    Log(const Log&) = delete;
    Log& operator=(const Log&) = delete;
};

#define LOG(level) \
if (level > LOG_CONFIG.reporting_level) , \
else Log().get(level) ← This macro will allow us to optimize our logging command and give us syntactic sugar

#endif
```

A Logging system

Source

src/engine/Log.cpp

```
#include "Log.h"

Log::Log() {
    file.open(GL_LOG_FILE, std::fstream::app);
}

Log::~Log()
{
    os << std::endl;
    file << os.str();
    printf(os.str().c_str());
    os.clear();
    file.close();
}

std::ofstream Log::file;

void Log::restart() {
    file.open(GL_LOG_FILE, std::fstream::trunc);
    file.close();
}

std::ostringstream& Log::get(LogLevel level) {
    if(!file) return os;

    // Log
    time_t now;
    struct tm * timeinfo;
    char date[19];

    time(&now);
    timeinfo = localtime (&now);
    strftime (date, 19, "%y-%m-%d %H:%M:%S", timeinfo);

    // Log
    os << date << " " << get_label(level) << ": \t";
    return os;
}

std::string Log::get_label(LogLevel type) {
    std::string label;
    switch(type) {
        case Debug:   label = "DEBUG"; break;
        case Info:    label = "INFO "; break;
        case Warning: label = "WARN "; break;
        case Error:   label = "ERROR"; break;
    }
    return label;
}
```

Log constructor opens file in append mode

Log destructor write the content of the buffer in the file and in the standard output (console)

Reset the file by opening in truncate mode and closing

Now replace all the cout and printf with proper logging.

Logging GLFW errors

GLFW triggers its own errors. Here is how to log them.

- Before the main function, define the LOG_CONFIG struct

```
LogConfig LOG_CONFIG = {};
```

- Configure the logging right after declaring the window :

```
LOG_CONFIG . reporting_level = Debug;
LOG_CONFIG . restart = true;
if(LOG_CONFIG . restart) {
    Log::restart();
}
```

- Then initialize GLFW with a proper logging. Set the GLFW error callback function

```
...
// Set viewport and clear color
glViewport(0, 0, SCREEN_WIDTH, SCREEN_HEIGHT);
glClearColor(0.0f, 0.0f, 0.2f, 1.0f);

if (glDebugMessageControlARB != nullptr) {
    glEnable(GL_DEBUG_OUTPUT_SYNCHRONOUS);
    glDebugMessageCallback((GLDEBUGPROCARB) debugGLErrorCallback, nullptr);
    GLuint unusedIds = 0;
    glDebugMessageControl(GL_DONT_CARE, GL_DONT_CARE, GL_DONT_CARE, 0, &unusedIds, GL_TRUE);
}

// Load
...
```

- debugGLErrorCallback will be defined in a new engine/GIDebug.hpp file :

```
#ifndef GL_DEBUG_HPP
#define GL_DEBUG_HPP

#include <GL/glew.h>
#include <algorithm>
#include <map>

#include "Log.h"

// Breakpoints that should ALWAYS trigger (EVEN IN RELEASE BUILDS) [x86]!
#ifndef _MSC_VER
#define eTB_CriticalBreakPoint() \
    if (!IsDebuggerPresent()) __debugbreak();
#else
#define eTB_CriticalBreakPoint() asm(" int $3");
#endif

const char *debugGISourceToStr(GLenum source) {
    static std::map<int, const char *> sources = {
        {GL_DEBUG_SOURCE_API, "API"}, 
        {GL_DEBUG_SOURCE_WINDOW_SYSTEM, "Window System"}, 
        {GL_DEBUG_SOURCE_SHADER_COMPILER, "Shader Compiler"}, 
        {GL_DEBUG_SOURCE_THIRD_PARTY, "Third Party"}, 
        {GL_DEBUG_SOURCE_APPLICATION, "Application User"}, 
        {GL_DEBUG_SOURCE_OTHER, "Other"}, 
    };
    return sources[source];
}

const char *debugGITypeToStr(GLenum type) {
    static std::map<int, const char *> types = {{GL_DEBUG_TYPE_ERROR, "Error"}, 
        {GL_DEBUG_TYPE_DEPRECATED_BEHAVIOR, "Deprecated Behaviour"}, 
        {GL_DEBUG_TYPE_UNDEFINED_BEHAVIOR, "Undefined Behaviour"}, 
        {GL_DEBUG_TYPE_PORTABILITY, "Portability"}, 
        {GL_DEBUG_TYPE_PERFORMANCE, "Performance"}, 
        {GL_DEBUG_TYPE_MARKER, "Marker"}, 
        {GL_DEBUG_TYPE_PUSH_GROUP, "Debug Push group"}, 
        {GL_DEBUG_TYPE_POP_GROUP, "Debug Pop Group"}, 
        {GL_DEBUG_TYPE_OTHER, "Other"}};
    return types[type];
}

const char *debugGISeverityToStr(GLenum severity) {
    static std::map<int, const char *> severities = {{GL_DEBUG_SEVERITY_HIGH, "High"}, 
        {GL_DEBUG_SEVERITY_MEDIUM, "Medium"}, 
        {GL_DEBUG_SEVERITY_LOW, "Low"}, 
        {GL_DEBUG_SEVERITY_NOTIFICATION, "Notification"}};
    return severities[severity];
}

void debugGLErrorCallback(GLenum source, GLenum type, GLuint id, GLenum severity, GLsizei length, const GLchar *message,
    GLvoid *userParam) {
    std::string glLog("- OpenGL -----");
    glLog.append("\n"); glLog.append("message: "); glLog.append(message);
    glLog.append("\n"); glLog.append("type: "); glLog.append(debugGITypeToStr(type));
    glLog.append("\n"); glLog.append("id: "); glLog.append(std::to_string(id));
    glLog.append("\n"); glLog.append("severity: "); glLog.append(debugGISeverityToStr(severity));

    if (type == GL_DEBUG_TYPE_ERROR) {
        LOG(Error) << glLog;
    } else if (type == GL_DEBUG_TYPE_DEPRECATED_BEHAVIOR || type == GL_DEBUG_TYPE_UNDEFINED_BEHAVIOR ||
        type == GL_DEBUG_TYPE_PORTABILITY || type == GL_DEBUG_TYPE_PERFORMANCE) {
        LOG(Warning) << glLog;
    } else {
        LOG(Debug) << glLog;
    }
}

#endif
```

Get graphics capabilities info

We want to get graphics capabilities. We will create a function that we will use just after GLEW init.

```
void logGIParams () {
    GLenum params[] = {
        GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS,
        GL_MAX_CUBE_MAP_TEXTURE_SIZE,
        GL_MAX_DRAW_BUFFERS,
        GL_MAX_FRAGMENT_UNIFORM_COMPONENTS,
        GL_MAX_TEXTURE_IMAGE_UNITS,
        GL_MAX_TEXTURE_SIZE,
        GL_MAX_VARYING_FLOATS,
        GL_MAX_VERTEX_ATTRIBS,
        GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS,
        GL_MAX_VERTEX_UNIFORM_COMPONENTS,
        GL_MAX_VIEWPORT_DIMS,
        GL_STEREO,
    };
    const char* names[] = {
        "GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS",
        "GL_MAX_CUBE_MAP_TEXTURE_SIZE",
        "GL_MAX_DRAW_BUFFERS",
        "GL_MAX_FRAGMENT_UNIFORM_COMPONENTS",
        "GL_MAX_TEXTURE_IMAGE_UNITS",
        "GL_MAX_TEXTURE_SIZE",
        "GL_MAX_VARYING_FLOATS",
        "GL_MAX_VERTEX_ATTRIBS",
        "GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS",
        "GL_MAX_VERTEX_UNIFORM_COMPONENTS",
        "GL_MAX_VIEWPORT_DIMS",
        "GL_STEREO",
    };
    LOG(Info) << "-----";
    LOG(Info) << "GL Context Params.";
    // integers - only works if the order is 0-10 integer return types
    for (int i = 0; i < 10; i++) {
        int v = 0;
        glGetIntegerv (params[i], &v);
        LOG(Info) << names[i] << " " << v;
    }
    // others
    int v[2];
    v[0] = v[1] = 0;
    glGetIntegerv (params[10], v);
    LOG(Info) << names[10] << " " << v[0] << " " << v[1];
    unsigned char s = 0;
    glGetBooleanv (params[11], &s);
    LOG(Info) << names[11] << " " << (unsigned int)s;
    LOG(Info) << "-----";
}
```

Use this function in Main.cpp just after the initialization and before the load part :

```
...
logGIParams();
// Load
...
```

Decipher graphics capabilities info

How to read our log result.

The log will output :

```
19-05-09 16:39:42 INFO : -----
19-05-09 16:39:42 INFO : GL Context Params:
19-05-09 16:39:42 INFO : GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS 192
19-05-09 16:39:42 INFO : GL_MAX_CUBE_MAP_TEXTURE_SIZE 16384
19-05-09 16:39:42 INFO : GL_MAX_DRAW_BUFFERS 8
19-05-09 16:39:42 INFO : GL_MAX_FRAGMENT_UNIFORM_COMPONENTS 4096
19-05-09 16:39:42 INFO : GL_MAX_TEXTURE_IMAGE_UNITS 32
19-05-09 16:39:42 INFO : GL_MAX_TEXTURE_SIZE 16384
19-05-09 16:39:42 INFO : GL_MAX_VARYING_FLOATS 64
19-05-09 16:39:42 INFO : GL_MAX_VERTEX_ATTRIBS 16
19-05-09 16:39:42 INFO : GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS 32
19-05-09 16:39:42 INFO : GL_MAX_VERTEX_UNIFORM_COMPONENTS 4096
19-05-09 16:39:42 INFO : GL_MAX_VIEWPORT_DIMS 16384 16384
19-05-09 16:39:42 INFO : GL_STEREO 0
19-05-09 16:39:42 INFO : -----
```

This tells me that my shader programmes can use 192 different textures each. I can access 96 different textures in the vertex shader, and 96 more in the fragment shader. My laptop can support only 8 textures, so if I want to write programmes that run nicely on both you can see that I would make sure that they don't use more than 8 textures at once. In theory my texture resolution can be up to 16384x16384, but multiply this by 4 bytes (for red+green+blue+alpha channels), and we see that 16k textures will use up my memory pretty quickly. I might get away with 8224x8224x4 bytes.

The "max uniform components" means that I can send 4096 floats to each shader. Each matrix is going to use 32 floats. So we can say that we have plenty of space there.

"Varying" floats are those sent from the vertex shader to the fragment shaders. Usually these are vectors, so we can say that we can send 16 4d vectors between shaders. Varyings are computationally expensive, and some devices still have a limit of 16 floats (4 vectors), so it's best to keep these to a minimum.

Vertex attributes are variables loaded from a mesh e.g. vertex points, texture coordinates, normals, per-vertex colours, etc. OpenGL means 4d vectors here. I would struggle to come up with more than about 6 useful per-vertex attributes, so no problem here. Draw buffers is useful for more advanced effects where we want to split the output from our rendering into different images - we can split this into 8 parts. And, sadly, my video card doesn't support stereo rendering.

GL state machine

In the first lesson, we told about the OpenGL State Machine. This means that once we set a state (like transparency, for example), it is then globally enabled for all future drawing operations, until we change it again. In GL parlance, setting a state is referred to as "binding" (for buffers of data), "enabling" (for rendering modes), or "using" for shader programmes.

The state machine can be very confusing. Lots of errors in OpenGL programmes come from setting a state by accident, forgetting to unset a state, or mixing up the numbering of different OpenGL indices. Some of the most useful state machine variables can be fetched during run-time. You probably don't need to write a function to log all of these states, but keep in mind that, if it all gets a bit confusing, you can check individual states.

A small framework

OOP refactoring : Window (1/2)

Now our OpenGL is initialized, we will refactor our code in a more OOP fashion.

We'll get most part of our initialization logic out of main.cpp to put it in a Window class, whose header will be :

```
#ifndef WINDOW_H
#define WINDOW_H

#include <GL/glew.h>
#include <SDL.h>
#include "Log.h"

#include <string>

class Window {
public:
    Window(const std::string &title);
    ~Window();

    bool init(int xPos, int yPos, int width, int height, bool isFullscreen);
    void logGLParams();
    void clearBuffer();
    void swapBuffer();
    void clean();

private:
    SDL_Window *window;
    SDL_GLContext context;
    const std::string &title;

    // Delete
    Window() = delete;
    Window(const Window &) = delete;
    Window &operator=(const Window &) = delete;
};

#endif
```

We can now send most of our initialization code into the Window functions.

OOP refactoring : Window (2/2)

Window.cpp

```
#include "Window.h"
#include "GIDebug.hpp"

#include <iostream>

using std::cout;
using std::endl;

Window::Window(const std::string &title)
: context(nullptr), title(title) {}

Window::~Window() {
    SDL_Quit();
    LOG(Info) << "Bye :)";
}

bool Window::init(int xPos, int yPos, int width, int height, bool isFullscreen) {
    int flags = SDL_WINDOW_OPENGL;
    if (isFullscreen) {
        flags = SDL_WINDOW_FULLSCREEN | SDL_WINDOW_OPENGL;
    }

    if (SDL_Init(SDL_INIT_EVERYTHING) != 0) {
        LOG(Error) << "SDL initialisation failed";
        LOG(Error) << SDL_GetError();
        return false;
    }

    LOG(Info) << "Subsystems initialised";

    // Sdl window
    window = SDL_CreateWindow(title.c_str(), xPos, yPos, width, height, flags);
    if (window) {
        LOG(Info) << "WindowSdl initialised";
    } else
        return false;

    // OpenGL context
    context = SDL_GL_CreateContext(window);
    if (context) {
        LOG(Info) << "OpenGL Context initialised";
    } else
        return false;

    // OpenGL setup
    glewExperimental = GL_TRUE;
    GLenum initGLEW(glewInit());
    if (initGLEW == GLEW_OK) {
        LOG(Info) << "GLEW initialised";
    } else
        return false;

    // Size of the viewport
    glViewport(0, 0, width, height);

    // Enable transparency
    glEnable(GL_BLEND);
    glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);

    // OpenGL error callback
    if (glDebugMessageControlARB != nullptr) {
        glEnable(GL_DEBUG_OUTPUT_SYNCHRONOUS);
        glDebugMessageCallback((GLDEBUGPROCARB) debugGLErrorCallback, nullptr);
        GLuint unusedIds = 0;
        glDebugMessageControl(GL_DONT_CARE, GL_DONT_CARE, GL_DONT_CARE, 0, &unusedIds, GL_TRUE);
    }

    // Window color
    glClearColor(0.0f, 0.0f, 0.2f, 1.0f);

    return true;
}

void Window::logGLParams() {
    GLenum params[] = {
        GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS,
        GL_MAX_CUBE_MAP_TEXTURE_SIZE,
        GL_MAX_DRAW_BUFFERS,
        GL_MAX_FRAGMENT_UNIFORM_COMPONENTS,
        GL_MAX_TEXTURE_IMAGE_UNITS,
        GL_MAX_TEXTURE_SIZE,
        GL_MAX_VARYING_FLOATS,
        GL_MAX_VERTEX_ATTRIBS,
        GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS,
        GL_MAX_VERTEX_UNIFORM_COMPONENTS,
        GL_MAX_VIEWPORT_DIMS,
        GL_STEREO,
    };
    const char* names[] = {
        "GL_MAX_COMBINED_TEXTURE_IMAGE_UNITS",
        "GL_MAX_CUBE_MAP_TEXTURE_SIZE",
        "GL_MAX_DRAW_BUFFERS",
        "GL_MAX_FRAGMENT_UNIFORM_COMPONENTS",
        "GL_MAX_TEXTURE_IMAGE_UNITS",
        "GL_MAX_TEXTURE_SIZE",
        "GL_MAX_VARYING_FLOATS",
        "GL_MAX_VERTEX_ATTRIBS",
        "GL_MAX_VERTEX_TEXTURE_IMAGE_UNITS",
        "GL_MAX_VERTEX_UNIFORM_COMPONENTS",
        "GL_MAX_VIEWPORT_DIMS",
        "GL_STEREO",
    };
    LOG(Info) << "-----";
    LOG(Info) << "GL Context Params:";
    // integers - only works if the order is 0-10 integer return types
    for (int i = 0; i < 10; i++) {
        int v = 0;
        glGetIntegerv(params[i], &v);
        LOG(Info) << names[i] << " " << v;
    }
    // others
    int v[2];
    v[0] = v[1] = 0;
    glGetIntegerv(params[10], v);
    LOG(Info) << names[10] << " " << v[0] << " " << v[1];
    unsigned char s = 0;
    glGetBooleanv(params[11], &s);
    LOG(Info) << names[11] << " " << (unsigned int)s;
    LOG(Info) << "";
}

void Window::clearBuffer() { glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT); }

void Window::swapBuffer() { SDL_GL_SwapWindow(window); }

void Window::clean() {
    SDL_DestroyWindow(window);
    SDL_GL_DeleteContext(context);
}
```

FPS counter (1/2)

We will add a FPS counter at the top of our window, which will be updated each quarter of seconds. First, we need to create a Timer class that will handle delta time calculation.

src/engine/Timer.h

```
#ifndef TIMER_H
#define TIMER_H

#include <SDL.h>

// Hold time related functions.
// In charge of computing the delta time and
// ensure smooth game ticking.
class Timer
{
public:
    Timer();
    virtual ~Timer();

    // Compute delta time as the number of milliseconds since last frame
    unsigned int computeDeltaTime();

    // Wait if the game run faster than the decided FPS
    void delayTime();

    // Get time since the game started, in seconds
    static double getTimeSinceStart() { return timeSinceStart; }

private:
    const static int FPS = 60;
    const static int frameDelay = 1000 / FPS;
    const unsigned int MAX_DT = 50;

    // Time in milliseconds when frame starts
    unsigned int frameStart;

    // Last frame start time in milliseconds
    unsigned int lastFrame;

    // Time it tooks to run the loop. Used to cap framerate.
    unsigned int frameTime;

    // Time in seconds since the game started
    static double timeSinceStart;
};

#endif
```

src/engine/Timer.cpp

```
#include "Timer.h"
#include <algorithm>

Timer::Timer() : frameStart(0), lastFrame(0), frameTime(0)
{
}

Timer::~Timer()
{
}

unsigned int Timer::computeDeltaTime()
{
    frameStart = SDL_GetTicks();
    unsigned int dt = frameStart - lastFrame;
    dt = std::min(dt, MAX_DT);
    timeSinceStart += static_cast<float>(dt) / 1000.0f;
    lastFrame = frameStart;
    return dt; // Clamp delta time
}

void Timer::delayTime()
{
    frameTime = SDL_GetTicks() - frameStart;
    if (frameTime < frameDelay) {
        SDL_Delay(frameDelay - frameTime);
    }
}

double Timer::timeSinceStart = 0.0f;
```

Note that `SDL_GetTicks` gives us a number of milliseconds. We will need to convert it into seconds later.

FPS counter (2/2)

Now we add a function that use the delta time to compute the FPS counter each quarter of a second.

Window.h

```
...
class Window {
public:
...
void updateFpsCounter(float dt);
...
private:
...
double previousSeconds;
double currentSeconds;
int frameCount;
...
}
```

Window.cpp

```
...
Window::Window(const std::string &title)
: context(nullptr), title(title), previousSeconds(0), currentSeconds(0), frameCount(0) {}

...

void Window::updateFpsCounter(float dt) {
    double elapsedSeconds;

    currentSeconds += dt;
    elapsedSeconds = currentSeconds - previousSeconds;
    /* limit text updates to 4 per second */
    if (elapsedSeconds > 0.25) {
        previousSeconds = currentSeconds;
        char tmp[128];
        double fps = (double)frameCount / elapsedSeconds;
        sprintf_s(tmp, "%s @ fps: %.2f", title.c_str(), fps);
        SDL_SetWindowTitle(window, tmp);
        frameCount = 0;
    }
    frameCount++;
}
```

We can compute the delta time in seconds and use it in the main :

Main.cpp

```
...
dt = static_cast<float>(timer.computeDeltaTime()) / 1000.0f;
window.updateFpsCounter(dt);
...
```

The game class

Even if our Window class will allow to tidy our code a lot, we will also add a Game class to handle the classic game loop. This will allow us to write code in a structured manner.

Game.h

```
#ifndef GAME_H
#define GAME_H

#include <SDL.h>
#include <GL/glew.h>
#include <vector>
#include <memory>

using std::vector;

// This game class runs a simple game loop
class Game
{
public:
    Game();
    virtual ~Game();

    void init(int screenWidth, int screenHeight);
    void load();
    void handleInputs();
    void update(float dt);
    void render();
    void clean();

    bool isRunning;
    int windowHeight, windowWidth;

private:
    GLuint shaderProgram;
    GLuint vao;
};
```

Game.cpp

```
#include "../engine/Game.h"

Game::Game() : isRunning(false), windowWidth(0), windowHeight(0) {}

Game::~Game() {}

void Game::init(int screenWidth, int screenHeight) {
    windowWidth = screenWidth;
    windowHeight = screenHeight;
    isRunning = true;
}

void Game::load() {}

void Game::handleInputs() {
    SDL_Event event;
    while (SDL_PollEvent(&event)) {
        switch (event.type) {
            case SDL_QUIT:
                isRunning = false;
                break;

            case SDL_KEYDOWN:
                if (event.key.keysym.sym == SDLK_ESCAPE) {
                    isRunning = false;
                }
                break;

            default:
                break;
        }
    }
}

void Game::update(float dt) {}

void Game::render() {}

void Game::clean() {}
```

The triangle with our framework

Now we can use our framework to draw the triangle.

Main.cpp

```
#include <SDL.h>
#include <iostream>

#include "engine/Game.h"
#include "engine/Timer.h"
#include "engine/Window.h"

using std::cout;
using std::endl;

LogConfig LOG_CONFIG = {};

int main(int argc = 0, char **argv = nullptr) {
    if (argc > 0) {
        for (int i = 0; i < argc; ++i) {
            cout << argv[i] << endl;
        }
    }

    const int SCREEN_WIDTH = 800;
    const int SCREEN_HEIGHT = 640;

    // Init logging
    LOG_CONFIG.reporting_level = Debug;
    LOG_CONFIG.restart = true;
    if (LOG_CONFIG.restart) {
        Log::restart();
    }

    // Init window
    Window window = Window("Hello SDL");
    if (!window.init(SDL_WINDOWPOS_CENTERED, SDL_WINDOWPOS_CENTERED, SCREEN_WIDTH, SCREEN_HEIGHT, false)) {
        return 1;
    }

    // Init game
    Timer timer;
    Game game;
    game.init(SCREEN_WIDTH, SCREEN_HEIGHT);
    game.load();

    // Delta time in seconds
    float dt;

    // Game loop
    while (game.isRunning) {
        dt = static_cast<float>(timer.computeDeltaTime()) / 1000.0f;
        window.updateFpsCounter(dt);

        game.handleInputs();
        game.update(dt);

        window.clearBuffer();
        game.render();
        window.swapBuffer();

        // Delay frame if game runs too fast
        timer.delayTime();
    }

    // Exit game
    game.clean();
    window.clean();
    return 0;
}
```

Game.cpp

```
#include "../engine/Game.h"

Game::Game() : isRunning(false), windowHeight(0), windowWidth(0) {}

Game::~Game() {}

void Game::init(int screenWidth, int screenHeight) {
    windowWidth = screenWidth;
    windowHeight = screenHeight;
    isRunning = true;
}

void Game::load() {
    GLfloat points[] = {0.0f, 0.5f, 0.0f, 0.5f, -0.5f, 0.0f, -0.5f, 0.0f};

    GLuint vbo = 0;
    glGenBuffers(1, &vbo);
    glBindBuffer(GL_ARRAY_BUFFER, vbo);
    glBufferData(GL_ARRAY_BUFFER, sizeof(points), points, GL_STATIC_DRAW);

    vao = 0;
    glGenVertexArrays(1, &vao);
    glBindVertexArray(vao);
    glEnableVertexAttribArray(0);
    glBindBuffer(GL_ARRAY_BUFFER, vbo);
    glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 0, NULL);

    const char* vertexShader =
        "#version 460\n"
        "in vec3 vp;"
        "void main () {"
        "    gl_Position = vec4 (vp, 1.0);"
        "}";

    const char* fragmentShader =
        "#version 460\n"
        "out vec4 color;"
        "void main () {"
        "    color = vec4 (0.8, 0.5, 0.0, 1.0);"
        "}";

    GLuint vs = glCreateShader(GL_VERTEX_SHADER);
    glShaderSource(vs, 1, &vertexShader, NULL);
    glCompileShader(vs);
    GLuint fs = glCreateShader(GL_FRAGMENT_SHADER);
    glShaderSource(fs, 1, &fragmentShader, NULL);
    glCompileShader(fs);

    shaderProgram = glCreateProgram();
    glAttachShader(shaderProgram, fs);
    glAttachShader(shaderProgram, vs);
    glLinkProgram(shaderProgram);
}

void Game::handleInputs() {
    SDL_Event event;
    while (SDL_PollEvent(&event)) {
        switch (event.type) {
            case SDL_QUIT:
                isRunning = false;
                break;

            case SDL_KEYDOWN:
                if (event.key.keysym.sym == SDLK_ESCAPE) {
                    isRunning = false;
                }
                break;

            default:
                break;
        }
    }
}

void Game::update(float dt) {}

void Game::render() {
    glUseProgram(shaderProgram);
    glBindVertexArray(vao);
    glDrawArrays(GL_TRIANGLES, 0, 3);
}

void Game::clean() {}
```

Shaders

Shaders

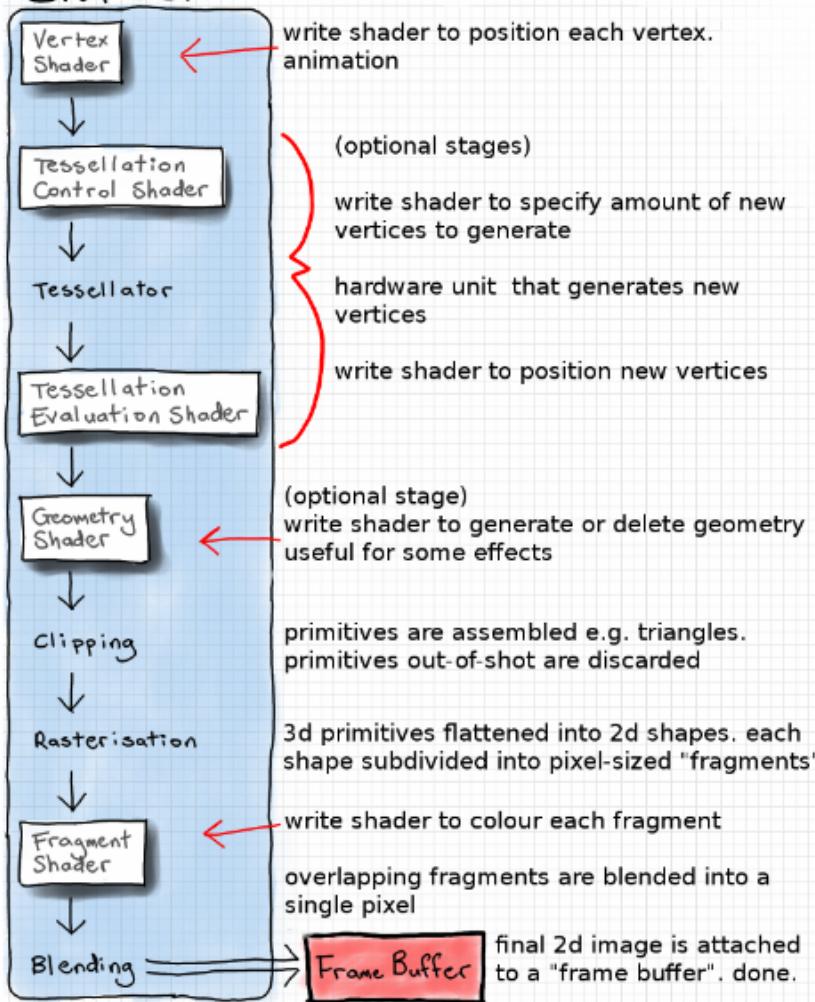
Shaders are the programmable part of a more general graphics pipeline. Graphics pipeline is the ensemble of steps the graphic card compute before sending pixels to display. It is a pipeline because inputs line to be computed one after another. Graphics card are massively parallel systems, so they compute several inputs at the same time (several thousands shaders at a time).

Here is the OpenGL 4 graphics pipeline. Steps in rectangles are customizable.

C.P.U.
glDrawArrays()
↓ go!

OpenGL 4 Hardware Pipeline

G.P.U.



A complete shader programme controls comprises a set of separate shader (mini-programmes) - one to control each stage. Each mini-programme - called a shader by OpenGL - is compiled, and the whole set are linked together to form the executable shader programme - called a program by OpenGL. You can see that the API differentiates functions into `glShader` and `glProgram`.

Each individual shader has a different job. At minimum, we usually have 1 vertex shader and 1 fragment shader per shader programme, but OpenGL 4 allows us to use some optional shaders too.

A shader is a small programme. OpenGL uses the GLSL, a C like programming language.

Fragment / Pixels

A pixel is a "picture element". In OpenGL lingo, pixels are the elements that make up the final 2d image that it draws inside a window on your display. A fragment is a pixel-sized area of a surface. A fragment shader determines the colour of each one. Sometimes surfaces overlap - we then have more than 1 fragment for 1 pixel. All of the fragments are drawn, even the hidden ones.

Each fragment is written into the framebuffer image that will be displayed as the final pixels. If depth testing is enabled it will paint the front-most fragments on top of the further-away fragments. In this case, when a farther-away fragment is drawn after a closer fragment, then the GPU is clever enough to skip drawing it, but it's actually quite tricky to organise the scene to take advantage of this, so we'll often end up executing huge numbers of redundant fragment shaders.

Vertex Shader

The vertex shader is responsible for transforming vertex positions into clip space, the final coordinate space that we must transform points to before OpenGL rasterises (flattens) our geometry into a 2d image. Vertex shaders can also be used to send data from the vertex buffer to fragment shaders. This vertex shader does nothing, except take in vertex positions that are already in clip space, and output them as final clip-space positions. We can write this into a plain text file called: test.vert.

```
#version 410

in vec3 vertex_position;

void main() {
    gl_Position = vec4(vertex_position, 1.0);
}
```

GLSL has some built-in data types that we can see here:

- `vec3` is a 3d vector that can be used to store positions, directions, or colours.
- `vec4` is the same but has a fourth component which, in this variable, is used to determine perspective. For now, we can leave it at 1.0, which means "don't calculate any perspective".

We can also see the `in` key-word for input to the programme from the previous stage. In this case the `vertex_position_local` is one of the vertex points from the object that we are drawing. GLSL also has an `out` keyword for sending a variable to the next stage.

The entry point to every shader is a `void main()` function.

The `gl_Position` variable is a built-in GLSL variable used to set the final clip-space position of each vertex.

The input to a vertex buffer (the `in` variables) are called per-vertex attributes, and come from blocks of memory on the graphics hardware memory called vertex buffers. We usually copy our vertex positions into vertex buffers before running our main loop. We will look at vertex buffers in the next chapter. This vertex shader will run one instance for every vertex in the vertex buffer.

Fragment Shader

Once all of the vertex shaders have computed the position of every vertex in clip space, then the fragment shader is run once for every pixel-sized space (fragment) between vertices. The fragment shader is responsible for setting the colour of each fragment. Write a new plain-text file: test.frag.

```
#version 410

uniform vec4 inputColour;
out vec4 fragColour;

void main() {
    fragColour = inputColour;
}
```

The uniform keyword says that we are sending in a variable to the shader programme from the CPU. This variable is global to all shaders within the programme, so we could also access it in the vertex shader if we wanted to.

The hardware pipeline knows that the first vec4 it gets as output from the fragment shader should be the colour of the fragment. The colours are rgba, or red, green, blue, alpha. The values of each component are floats between 0.0 and 1.0, (not between 0 and 255). The alpha channel output can be used for a variety of effects, which you define by setting a blend mode in OpenGL. It is commonly used to indicate opacity (for transparent effects), but by default it does nothing.

Some OpenGL shader functions

For a complete list of OpenGL shader functions see the Quick Reference Card <http://www.opengl.org/documentation/glsl/>.

The most useful functions are tabulated below. We will implement all of these:

- `glCreateShader()` - create a variable for storing a shader's code in OpenGL. returns GLuint index to it.
- `glShaderSource()` - copy shader code from C string into an OpenGL shader variable
- `glCompileShader()` - compile an OpenGL shader variable that has code in it
- `glGetShaderiv()` - can be used to check if compile found errors
- `glGetShaderInfoLog()` - creates a string with any error information
- `glDeleteShader()` - free memory used by an OpenGL shader variable
- `glCreateProgram()` - create a variable for storing a combined shader programme in OpenGL. returns GLuint index to it.
- `glAttachShader()` - attach a compiled OpenGL shader variable to a shader programme variable
- `glLinkProgram()` - after all shaders are attached, link the parts into a complete shader programme
- `glValidateProgram()` - check if a program is ready to execute. information stored in a log
- `glGetProgramiv()` - can be used to check for link and validate errors
- `glGetProgramInfoLog()` - writes any information from link and validate to a C string
- `glUseProgram()` - switch to drawing with a specified shader programme
- `glGetActiveAttrib()` - get details of a numbered per-vertex attribute used in the shader
- `glGetAttribLocation()` - get the unique "location" identifier of a named per-vertex attribute
- `glGetUniformLocation()` - get the unique "location" identifier of a named uniform variable
- `glGetActiveUniform()` - get details of a named uniform variable used in the shader
- `glUniform{1234}{fd}()` - set the value of a uniform variable of a given shader (function name varies by dimensionality and data type)
- `glUniform{1234}{fd}v()` - same as above, but with a whole array of values
- `glUniformMatrix{234}{fd}v()` - same as above, but for matrices of dimensions 2x2,3x3, or 4x4

Logging shader and linking errors

Shader error

After the vertex shader compilation, add this code :

```
int params = -1;
glGetShaderiv(vs, GL_COMPILE_STATUS, &params);
if (params != GL_TRUE)
{
    LOG(Error) << "GL vertex shader index " << vs << " did not compile.";
    printShaderInfoLog(vs);
    return -1;
}
```

With :

```
void printShaderInfoLog(GLuint shader_index)
{
    int maxLength = 2048;
    int actualLength = 0;
    char log[2048];
    glGetShaderInfoLog(shaderIndex, maxLength, &actualLength, log);
    LOG(Info) << "Shader info log for GL index" << shaderIndex;
    LOG(Info) << log;
}
```

Do the same for the fragment shader.

Linking error

After the programme compilation, add this code :

```
params = -1;
glGetProgramiv(shaderProgram, GL_LINK_STATUS, &params);
if (params != GL_TRUE)
{
    LOG(Error) << "Could not link shader programme GL index " << shaderProgram;
    printProgrammeInfoLog(shaderProgram);
}
```

With :

```
void printProgrammeInfoLog(GLuint programId)
{
    int maxLength = 2048;
    int actualLength = 0;
    char log[2048];
    glGetProgramInfoLog(id, maxLength, &actualLength, log);
    LOG(Info) << "program info log for GL index" << programId;
    LOG(Info) << log;
}
```

Logging usual errors

To custom our shader output, we can pass variables from the application (run by the CPU) to the shaders programmes. Those variable will keep the same value for all shaders computations, that is why they are called uniforms.

One of the more common errors is mixing up the "location" of uniform variables. Another is where an attribute or uniform variable is not "active"; not actually used in the code of the shader, and has been optimised out by the shader compiler. We can check this by printing it, and we can print all sorts of other information as well. Here, programId is the index of the shader programme.

```
void printAllParams(GLuint programId)
{
    LOG(Info) << "-----";
    LOG(Info) << "Shader programme " << programId << " info:";
    int params = -1;
    glGetProgramiv(programId, GL_LINK_STATUS, &params);
    LOG(Info) << "GL_LINK_STATUS = " << params;

    glGetProgramiv(programId, GL_ATTACHED_SHADERS, &params);
    LOG(Info) << "GL_ATTACHED_SHADERS = " << params;

    glGetProgramiv(programId, GL_ACTIVE_ATTRIBUTES, &params);
    LOG(Info) << "GL_ACTIVE_ATTRIBUTES = " << params;
    for (GLuint i = 0; i < (GLuint)params; i++)
    {
        char name[64];
        int maxLength = 64;
        int actualLength = 0;
        int size = 0;
        GLenum type;
        glGetActiveAttrib(programId, i, maxLength, &actualLength, &size, &type, name);
        if (size > 1)
        {
            for (int j = 0; j < size; j++)
            {
                char longName[77];
#if __linux__
                sprintf(longName, "%s[%i]", name, j);
#else
                sprintf_s(longName, "%s[%i]", name, j);
#endif
                int location = glGetUniformLocation(programId, longName);
                LOG(Info) << " " << i << ") type:" << GLTypeToString(type) << " name:" << longName << " location:" << location;
            }
        }
        else
        {
            int location = glGetUniformLocation(programId, name);
            LOG(Info) << " " << i << ") type:" << GLTypeToString(type) << " name:" << name << " location:" << location;
        }
    }

    glGetProgramiv(programId, GL_ACTIVE_UNIFORMS, &params);
    LOG(Info) << "GL_ACTIVE_UNIFORMS = " << params;
    for (GLuint i = 0; i < (GLuint)params; i++)
    {
        char name[64];
        int maxLength = 64;
        int actualLength = 0;
        int size = 0;
        GLenum type;
        glGetActiveUniform(programId, i, maxLength, &actualLength, &size, &type, name);
        if (size > 1)
        {
            for (int j = 0; j < size; j++)
            {
                char longName[77];
                sprintf(longName, "%s[%i]", name, j);
                int location = glGetUniformLocation(programId, longName);
                LOG(Info) << " " << i << ") type:" << GLTypeToString(type) << " name:" << longName << " location:" << location;
            }
        }
        else
        {
            int location = glGetUniformLocation(programId, name);
            LOG(Info) << " " << i << ") type:" << GLTypeToString(type) << " name:" << name << " location:" << location;
        }
    }
    printProgrammeInfoLog(programId);
}
```

The interesting thing here are the printing of the attribute and uniform "locations". Sometimes uniforms or attributes are themselves arrays of variables - when this happens size is > 1, and we loop through and print each index' location separately.

Here is a home-made function for printing the GL data type as a string (normally it is an enum which doesn't look very meaningful when printed as an integer) :

```
const char *GLTypeToString(GLenum type)
{
    switch (type)
    {
        case GL_BOOL:
            return "bool";
        case GL_INT:
            return "int";
        case GL_FLOAT:
            return "float";
        case GL_FLOAT_VEC2:
            return "vec2";
        case GL_FLOAT_VEC3:
            return "vec3";
        case GL_FLOAT_VEC4:
            return "vec4";
        case GL_FLOAT_MAT2:
            return "mat2";
        case GL_FLOAT_MAT3:
            return "mat3";
        case GL_FLOAT_MAT4:
            return "mat4";
        case GL_SAMPLER_2D:
            return "sampler2D";
        case GL_SAMPLER_3D:
            return "sampler3D";
        case GL_SAMPLER_CUBE:
            return "samplerCube";
        case GL_SAMPLER_2D_SHADOW:
            return "sampler2DShadow";
        default:
            return "other";
    }
}
```

Program validation

You can also "validate" a shader programme before using it. Only do this during development, because it is quite computationally expensive. When a programme is not valid, the details will be written to the program info log.

```
bool isValid(GLuint programId)
{
    glValidateProgram(programId);
    int params = -1;
    glGetProgramiv(programId, GL_VALIDATE_STATUS, &params);
    LOG(Info) << "program " << programId << " GL_VALIDATE_STATUS = " << params;
    if (params != GL_TRUE)
    {
        printProgrammeInfoLog(programId);
        return false;
    }
    return true;
}
```

OOP refactorisation (1/3)

Now create a Shader class header and source file, to store our shader functions. Header can be :

Shader.h

```
#ifndef SHADER_H
#define SHADER_H

#include "Log.h"
#include <GL/glew.h>

class Shader
{
public:
    Shader();
    virtual ~Shader();

    GLuint programId;

    void compileVertexShader();
    void compileFragmentShader();
    void createShaderProgram();

    // Sets the current shader as active
    Shader& use();

private:
    GLuint vs;
    GLuint fs;

    const char* vertexShader =
        "#version 410\n"
        "in vec3 vp;"
        "void main () {"
        "    gl_Position = vec4 (vp, 1.0);"
        "}";

    const char* fragmentShader =
        "#version 410\n"
        "out vec4 frag_colour;"
        "void main () {"
        "    frag_colour = vec4 (0.8, 0.5, 0, 1.0);"
        "}";

    void checkShaderErrors(GLuint shader, std::string shaderType);

    void printShaderInfoLog(GLuint shaderIndex);
    void printProgramInfoLog(GLuint programId);
    const char *GLTypeToString(GLenum type);
    void printAllParams(GLuint programId);
    bool isValid(GLuint programId);

    Shader(const Shader &) = delete;
    Shader& operator=(const Shader &) = delete;
};

#endif
```

OOP refatorisation (2/3)

Shader.cpp

```
#include "Shader.h"

Shader::Shader() {}

Shader::~Shader() {}

void Shader::compileVertexShader()
{
    vs = glCreateShader(GL_VERTEX_SHADER);
    glShaderSource(vs, 1, &vertexShader, NULL);
    glCompileShader(vs);
    checkShaderErrors(vs, "vertex");
}

void Shader::compileFragmentShader()
{
    fs = glCreateShader(GL_FRAGMENT_SHADER);
    glShaderSource(fs, 1, &fragmentShader, NULL);
    glCompileShader(fs);
    checkShaderErrors(fs, "fragment");
}

void Shader::createShaderProgram()
{
    // Create program
    programId = glCreateProgram();
    glAttachShader(programId, fs);
    glAttachShader(programId, vs);
    glLinkProgram(programId);

    // Check for linking error
    int params = -1;
    glGetProgramiv(programId, GL_LINK_STATUS, &params);
    if (params != GL_TRUE)
    {
        LOG(Error) << "Could not link shader programme GL index " << programId;
        printProgramInfoLog(programId);
    }
    if (!isValid(programId))
    {
        LOG(Error) << "Could not validate shader" << programId;
    }

    // Delete shaders for they are no longer used
    glDeleteShader(vs);
    glDeleteShader(fs);
}

Shader &Shader::use()
{
    glUseProgram(programId);
    return *this;
}

void Shader::checkShaderErrors(GLuint shader, std::string shaderType)
{
    int params = -1;
    glGetShaderiv(shader, GL_COMPILE_STATUS, &params);
    if (params != GL_TRUE)
    {
        LOG(Error) << "GL " << shaderType << " shader index " << shader << " did not compile.";
        printShaderInfoLog(shader);
    }
}

void Shader::printShaderInfoLog(GLuint shaderIndex)
{
    int maxLength = 2048;
    int actualLength = 0;
    char log[2048];
    glGetShaderInfoLog(shaderIndex, maxLength, &actualLength, log);
    LOG(Info) << "Shader info log for GL index" << shaderIndex;
    LOG(Info) << log;
}

void Shader::printProgramInfoLog(GLuint programId)
{
    int maxLength = 2048;
    int actualLength = 0;
    char log[2048];
    glGetProgramInfoLog(programId, maxLength, &actualLength, log);
    LOG(Info) << "program info log for GL index" << programId;
    LOG(Info) << log;
}

const char *Shader::GLTypeToString(GLenum type)
{
    switch (type)
    {
        case GL_BOOL:
            return "bool";
        case GL_INT:
            return "int";
        case GL_FLOAT:
            return "float";
        case GL_FLOAT_VEC2:
            return "vec2";
        case GL_FLOAT_VEC3:
            return "vec3";
        case GL_FLOAT_VEC4:
            return "vec4";
        case GL_FLOAT_MAT2:
            return "mat2";
        case GL_FLOAT_MAT3:
            return "mat3";
        case GL_FLOAT_MAT4:
            return "mat4";
        case GL_SAMPLER_2D:
            return "sampler2D";
        case GL_SAMPLER_3D:
            return "sampler3D";
        case GL_SAMPLER_CUBE:
            return "samplerCube";
        case GL_SAMPLER_2D_SHADOW:
            return "sampler2DShadow";
        default:
            return "other";
    }
}

void Shader::printAllParams(GLuint programId)
{
    LOG(Info) << "-----";
    LOG(Info) << "Shader programme " << programId << " info:";

    int params = -1;
    glGetProgramiv(programId, GL_LINK_STATUS, &params);
    LOG(Info) << "GL_LINK_STATUS = " << params;

    glGetProgramiv(programId, GL_ATTACHED_SHADERS, &params);
    LOG(Info) << "GL_ATTACHED_SHADERS = " << params;

    glGetProgramiv(programId, GL_ACTIVE_ATTRIBUTES, &params);
    LOG(Info) << "GL_ACTIVE_ATTRIBUTES = " << params;
    for (GLuint i = 0; i < (GLuint)params; i++)
    {
        char name[64];
        int maxLength = 64;
        int actualLength = 0;
        int size = 0;
        GLenum type;
        glGetActiveAttrib(programId, i, maxLength, &actualLength, &size, &type, name);
        if (size > 1)
        {
            for (int j = 0; j < size; j++)
            {
                char longName[77];
                #if __linux__
                    sprintf(longName, "%s[%i]", name, j);
                #else
                    sprintf_s(longName, "%s[%i]", name, j);
                #endif
                int location = glGetUniformLocation(programId, longName);
                LOG(Info) << " " << i << " type:" << GLTypeToString(type) << " name:" << longName << " location:" << location;
            }
        }
        else
        {
            int location = glGetUniformLocation(programId, name);
            LOG(Info) << " " << i << " type:" << GLTypeToString(type) << " name:" << name << " location:" << location;
        }
    }

    glGetProgramiv(programId, GL_ACTIVE_UNIFORMS, &params);
    LOG(Info) << "GL_ACTIVE_UNIFORMS = " << params;
    for (GLuint i = 0; i < (GLuint)params; i++)
    {
        char name[64];
        int maxLength = 64;
        int actualLength = 0;
        int size = 0;
        GLenum type;
        glGetActiveUniform(programId, i, maxLength, &actualLength, &size, &type, name);
        if (size > 1)
        {
            for (int j = 0; j < size; j++)
            {
                char longName[77];
                sprintf(longName, "%s[%i]", name, j);
                int location = glGetUniformLocation(programId, longName);
                LOG(Info) << " " << i << " type:" << GLTypeToString(type) << " name:" << longName << " location:" << location;
            }
        }
        else
        {
            int location = glGetUniformLocation(programId, name);
            LOG(Info) << " " << i << " type:" << GLTypeToString(type) << " name:" << name << " location:" << location;
        }
    }
    printProgramInfoLog(programId);
}

bool Shader::isValid(GLuint programId)
{
    glValidateProgram(programId);
    int params = -1;
    glGetProgramiv(programId, GL_VALIDATE_STATUS, &params);
    LOG(Info) << "program " << programId << " GL_VALIDATE_STATUS = " << params;
    if (params != GL_TRUE)
    {
        printProgramInfoLog(programId);
        return false;
    }
    return true;
}
```

OOP refactorisation (3/3)

We can now use the shader class into the game class. First, replace the program id fo the game's header by a Shader variable, called "shader". Then :

Game.cpp

```
#include "../engine/Game.h"

Game::Game() : isRunning(false), windowHeight(0), windowWidth(0) {}

Game::~Game() {}

void Game::init(int screenWidth, int screenHeight) {
    windowWidth = screenWidth;
    windowHeight = screenHeight;
    isRunning = true;
}

void Game::load() {
    GLfloat points[] = {0.0f, 0.5f, 0.0f, 0.5f, -0.5f, 0.0f, -0.5f, -0.5f, 0.0f};

    GLuint vbo = 0;
    glGenBuffers(1, &vbo);
    glBindBuffer(GL_ARRAY_BUFFER, vbo);
    glBufferData(GL_ARRAY_BUFFER, sizeof(points), points, GL_STATIC_DRAW);

    vao = 0;
    glGenVertexArrays(1, &vao);
    glBindVertexArray(vao);
    glEnableVertexAttribArray(0);
    glBindBuffer(GL_ARRAY_BUFFER, vbo);
    glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 0, NULL);

    shader.compileVertexShader();
    shader.compileFragmentShader();
    shader.createShaderProgram();
}

void Game::handleInputs() {
    SDL_Event event;
    while (SDL_PollEvent(&event)) {
        switch (event.type) {
            case SDL_QUIT:
                isRunning = false;
                break;

            case SDL_KEYDOWN:
                if (event.key.keysym.sym == SDLK_ESCAPE) {
                    isRunning = false;
                }
                break;

            default:
                break;
        }
    }
}

void Game::update(float dt) {}

void Game::render() {
    shader.use();
    glBindVertexArray(vao);
    glDrawArrays(GL_TRIANGLES, 0, 3);
}

void Game::clean() {}
```

Best practices

- All uniform variables are initialised to 0 when a programme links, so you only need to initialise them if the initial value should be something else. Example: you might want to set matrices to the identity matrix, rather than a zeroed matrix.
- Calling `glUniform` is quite expensive during run-time. Structure your programme so that `glUniform` is only called when the value needs to change. This might be the case every time that you draw a new object (e.g. its position might be different), but some uniforms may not change often (e.g. projection matrix).
- Calling `glGetUniformLocation` during run-time can be expensive. It is best to do this during initialisation of the component that updates the uniform e.g. the virtual camera for the projection and view matrices, or a renderable object in the scene for the model matrix. Store the uniform locations once, then call `glUniform` as needed, rather than updating everything every frame.
- When calling `glGetUniformLocation`, it returns -1 if the uniform variable wasn't found to be active. You can check for this. Usually it means that either you've made a typo in the name, or the variable isn't actually used anywhere in the shader, and has been "optimised out" by the compiler/linker.
- Modifying attributes (vertex buffers) during run-time is extremely expensive. Avoid.
- Get your shaders to do as much work as is possible; because of their parallel nature they are much faster than looping on the CPU for most tasks.
- Drawing lots of separate, small objects at once does not make efficient use of the GPU, as most parallel shader slots will be empty, and separate objects must be drawn in series. Where possible, merge many, smaller objects into fewer, larger objects.

Extending your shader class

- Create a "reload my shaders" function. Bind this to a keyboard key. If you get this working your should be able to live edit your shaders and see what they do without restarting your programme. This should speed you up.
- Larger projects will use many different shader programmes. It would make sense to have a Shader Manager interface or class to load shaders, and to make sure that shaders are re-used, rather than loaded multiple times.
- In the future we will sometimes use geometry and tessellation shaders, so we can consider upgrading our shader functions to handle more than just vertex and fragment shaders. We can just set some boolean flags to true or false to indicate which types of shader have been loaded before attaching and linking.
- If you have a shader manager, and have written a function along the lines of `setUniform (shaderIndex, value)` from the manager, it could then check to make sure that shader is in use first (a common cause of error).

Vertex Buffer Objects

Why VBO ?

A vertex buffer object (VBO) is nothing fancy - it's just an array of data (usually floats). We already had a look at the most basic use of vertex buffer objects in the Triangle tutorial. We know that we can describe a mesh in an array of floats; {x,y,z,x,y,z..x,y,z}. And we also know that we need to use a Vertex Array Object to tell OpenGL that the array is divided into variables of 3 floats each.

The key idea of VBOs is this: in the old, "immediate-mode" days of OpenGL, before VBOs, we would define the vertex data in main memory (RAM), and copy them one-by-one each time that we draw. With VBOs, we copy the whole lot into a buffer before drawing starts, and this sits on the graphics hardware memory instead. This is much more efficient for drawing because, although the bus between the CPU and the GPU is very wide, a bottle-neck for drawing performance is created by stalling drawing operations to send OpenGL commands from the CPU to the GPU. To avoid this we try to keep as much data and processing on the graphics hardware as we can.

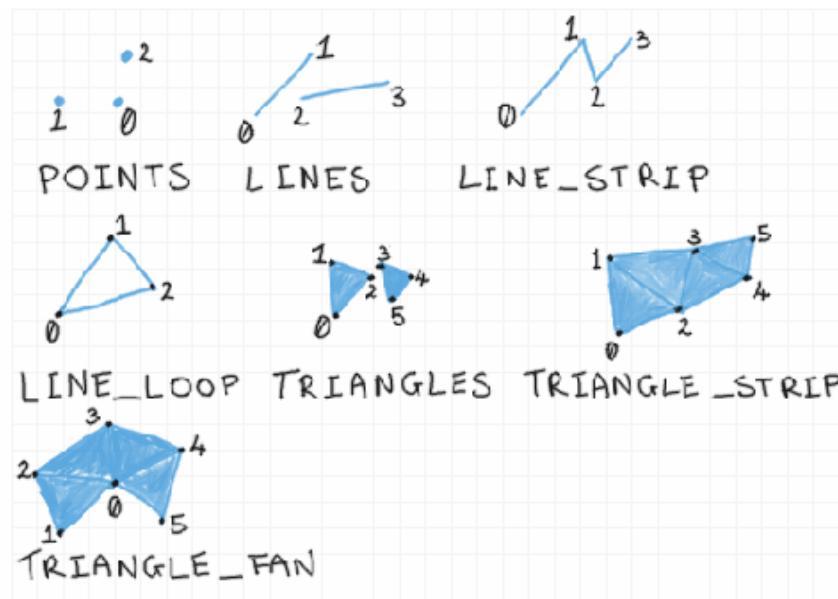
A VBO is not an "object" in the object-oriented programming sense, it is a simple array of data. OpenGL has several other types of data that it refers to as singular "buffer objects". The term "object" here implies that the GL gives us a handle or identifier number to the whole buffer to interact with, rather than a traditional address to the first element in the buffer.

We will look at managing vertex buffers in a little bit more depth. We will need this for enabling lighting and texturing effects later. We will break down the interface, and visualise how interpolation between the vertex shader and the fragment shader works.

Rendering Different Primitive Types

Drawing modes

We used `glDrawArrays(GL_TRIANGLES, 0, 3)` to draw a triangle. But there are several other drawing modes :



Try changing your `GL_TRIANGLES` parameter. You can see that points and lines are going to be useful for drawing things like charts or outlines. Triangle strip is a slightly more efficient method for drawing ribbon-like shapes. You can actually change the size of the points in the vertex shader, so they are quite versatile.

Wire-frame

It's much easier to use the `glPolygonMode` built-in function than to attempt to draw everything with `GL_LINES`. Call `glPolygonMode(GL_FRONT_AND_BACK, GL_LINE);` before rendering.

Using Multiple Vertex Buffers for One Object

We can store more than just 3d points in a vertex buffer. Common uses also include :

- 2d texture coordinates that describe how to fit an image to a surface
- 3d normals that describe which way a surface is facing so that we can calculate lighting.

So it's quite likely that most of your objects will have 2 or 3 vertex buffers each.

You can use vertex buffers to hold any data that you like; the key thing is that the data is retrieved once per vertex. If you tell OpenGL to draw an array of 3 points, using a given vertex array object, then it is going to launch 3 vertex shaders in parallel, and each vertex shader will get a one variable from each of the attached arrays; the first vertex shader will get the first 3d point, 2d texture coordinate, and 3d normal, the second vertex shader will get the second 3d point, 2d texture coordinate, and 3d normal, and so on, where the number of variables, and size of each variable is laid out in the vertex array object.

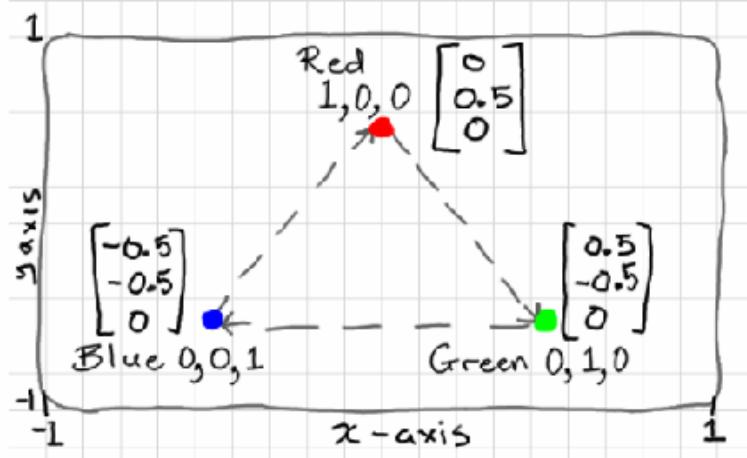
Colour interpolation on vertices 1/3

Data

Vertex colors are very seldom used in practise, but most modern GL tutorials will get you to create a buffer of colours as a second vertex buffer. Why? Because it's easy to visualise how interpolation works with colors.

Game.cpp

```
GLfloat points[] = {  
    0.0f, 0.5f, 0.0f,  
    0.5f, -0.5f, 0.0f,  
    -0.5f, -0.5f, 0.0f  
};  
  
GLfloat colors[] = {  
    1.0f, 0.0f, 0.0f,  
    0.0f, 1.0f, 0.0f,  
    0.0f, 0.0f, 1.0f  
};
```



VBO

Now for each one we can create a GL vertex buffer object, bind it in the state machine, and copy the array of values into it. Both colous, and points have 9 components each (3 vertices with 3 components per vertex). We set up one buffer after the other, because the state machine can only have 1 buffer bound at a time.

Game.cpp

```
GLuint pointsVbo = 0;  
glGenBuffers (1, &pointsVbo);  
glBindBuffer (GL_ARRAY_BUFFER, pointsVbo);  
glBufferData (GL_ARRAY_BUFFER, sizeof (points), points, GL_STATIC_DRAW);  
  
GLuint colorsVbo = 0;  
glGenBuffers (1, &colorsVbo);  
glBindBuffer (GL_ARRAY_BUFFER, colorsVbo);  
glBufferData (GL_ARRAY_BUFFER, sizeof (colors), colors, GL_STATIC_DRAW);
```

Colour interpolation on vertices 2/3

VAO

Our object now consists of 2 vertex buffers, which will be input "attribute" variables to our vertex shader. We set up the layout of both of these with a single vertex attribute object - the VAO represents our complete object, so we no longer need to keep track of the individual VBOs.

Game.cpp

```
GLuint vao = 0;
glGenVertexArrays(1, &vao);
 glBindVertexArray(vao);
 glBindBuffer(GL_ARRAY_BUFFER, pointsVbo);
 glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 0, NULL);
 glBindBuffer(GL_ARRAY_BUFFER, coloursVbo);
 glVertexAttribPointer(1, 3, GL_FLOAT, GL_FALSE, 0, NULL);
```

Note that we have to bind each VBO before calling `glVertexAttribPointer()`, which describes the layout of each buffer to the VAO.

The first parameter of `glVertexAttribPointer()` asks for an index. This is going to map to the indices in our vertex shader, so we need to give each attribute here a unique index. I will give my points index 0, and the colours index 1. We will make these match up to the variables in our vertex shader later. If you accidentally leave both indices at 0 (easy enough to do when copy-pasting code), then your colors will be read from the position values so $x \rightarrow$ red $y \rightarrow$ green and $z \rightarrow$ blue.

Both buffers contain arrays of floating point values, hence `GL_FLOAT`, and each variable has 3 components each, hence the 3 in the second parameter. If you accidentally get this parameter wrong (quite a common mistake), then the vertex shaders will be given variables made from the wrong components (e.g. position x, y, z gets values read from a, y, z, x).

That is the mesh side taken care of - you only need to do this part once, when creating the object. There is no need to repeat this code inside the rendering loop. Just keep track of the VAO index for each type of mesh that you create.

Enable Vertex Arrays 0 and 1 in a VAO

Attributes are disabled by default in OpenGL 4. We need to explicitly enable them too. This is easy to get wrong or overlook. We use a function called `glEnableVertexAttribArray()` to enable each one. This function only affects the currently bound vertex array object. This means that when we do this now, it will only affect our attributes, above. We will need to bind every new vertex array and repeat this procedure for those too.

Game.cpp

```
glEnableVertexAttribArray(0);
glEnableVertexAttribArray(1);
```

NB : Other GL incarnations don't have vertex attribute objects, and need to explicitly enable and disable the attributes every time a new type of object is drawn, which changes the enabled attributes globally in the state machine. We don't need to worry about that in OpenGL 4; the vertex attribute object will remember its enabled attributes, i.e. they are no longer global states. This isn't clear in the official documentation, nor is it explained properly in other tutorials - it's very easy to get a false-positive misunderstanding of how these things work and run into hair-pulling problems later when you're trying to draw 2 different shapes.

Colour interpolation on vertices 3/3

Change the vertex shader

We want our shader to render using the second vertex buffer, so we need to add a second attribute to the top of the vertex shader. We are going to add the OpenGL 4 layout prefix to each attribute. This lets us manually specify a location for each attribute - and we are going to match this up to the index that we gave each one in `glVertexAttribPointer`. If you don't specify a location, the shader programme linker will automatically assign one. You can query this, which we did in Shaders, but you would usually specify it manually so you can see the values as you write.

Shader.h

```
const char* vertexShader =
"#version 430\n"

"layout(location = 0) in vec3 vertexPosition;"
"layout(location = 1) in vec3 vertexColor;"

"out vec3 color;"

"void main ()"
"{"
"color = vertexColor;"
"gl_Position = vec4 (vertexPosition, 1.0);"
"}";
```

Now, the other interesting change here is that I took the `vertex_colour` input attribute and gave it to an output variable, which I called `colour`. This is where it gets interesting. It has no effect on our vertex shader, but outputs it to the next stage in the programmable hardware pipeline. In our case - the fragment shader.

Change the fragment shader

We can rewrite the fragment shader to use our new colour variable as an input, which will colour each fragment directly. Note that we add an `in` prefix to retrieve a variable from a previous stage. This in/out convention is a little different to how other OpenGL versions work. Our output fragment colour needs to be `r,g,b,a` (4 components) so we can just add a `1` to the end of our 3-component colour by casting it as a `vec4`.

Shader.h

```
const char *fragment_shader =
"#version 430\n"

"in vec3 color;"
"out vec4 fragColor;"

"void main ()"
"{"
"fragColor = vec4 (color, 1.0);"
"}";
```

Vertex Shader to Fragment Shader Interpolation

Now, remember, our triangle has only 3 vertices, but 1 fragment for every pixel-sized area of the surface. This means that we have 3 colour outputs from vertex shaders, and perhaps 100 colour inputs to each fragment shader. How does this work?

The answer is that each fragment shader gets an interpolated colour based on its position on the surface. The fragment exactly on the red corner of the triangle will be completely red (1.0, 0.0, 0.0). A fragment exactly half-way between the blue and red vertices, along the edge of the triangle, will be purple; half red, half blue, and no green: (0.5, 0.0, 0.5). A fragment exactly in the middle of the triangle will be an equal mixture of all 3 colours; (0.3333, 0.3333, 0.333).

Keep in mind that this will happen with any other vertex buffer attributes that you send to the fragment shader; normals will be interpolated, and texture coordinates will be interpolated to each fragment. This is really handy, and we will exploit it for lots of interesting per-pixel effects. But it's quite common to misunderstand this when getting started with shaders - vertex shader outputs are not sending a constant variable from a vertex shader to all fragment shaders. We use uniform variables for that.

"Winding" and Back-Face Culling

The last thing that you should know about is a built-in rendering optimisation called back-face culling. This gives a hint to GL so that it can throw away the hidden "back" or inside faces of a mesh. This should remove half of the vertex shaders, and half of the fragment shader instances from the GPU - allowing you to render things twice as large in the same time. It's not appropriate all of the time - you might want our 2d triangle to spin and show both sides.

The only things that you need specify are if clock-wise vertex "winding" order means the front, or the back of each face, and set the GL state to enable culling of that side. Our triangle, you can see, is given in clock-wise order. Most mesh formats actually go the other way, so it pays to test this before wondering why a mesh isn't showing up at all!

Window.cpp

```
...  
    // Enable cull face optimization  
    glEnable(GL_CULL_FACE);      // Cull face  
    glCullFace(GL_BACK);        // Cull back face  
    glFrontFace(GL_CW);         // GL_CCW for Counter Clock-Wise  
...
```

Try switching to counter clock-wise to make sure that the triangle disappears. If you were to rotate it around now you'd see the other side was visible. As with other GL states, this culling is enabled globally, in the state machine, you can enable and disable it between calls to `glDrawArrays` so that some objects are double-sided, and some are single-sided, etc. Keep in mind which winding order you are making new shapes in.

Using matrices with OpenGL

Row Order and Column Order

When you are computing vectors and matrices in C code, you need to stick to a layout convention. The two layouts are row major and column major. You can use either convention, but it needs to be consistent. OpenGL people tend to favour column-major matrices, and Direct3D people row-major. You can swap between row-major and column-major by transposing a matrix. This flips it along the main diagonal (top-left to bottom-right diagonal). It's easy to spot if a 4X4 matrix is in row order because the transformation components will be on the bottom-row, and there will be zeros in the right-most column (as in the example drawn above). Column-major matrices will have the transformation in the right-most column.

The layout convention affects the order of multiplication, as detailed in below:

$$\begin{bmatrix} 1 & 0 & 0 & T_x \\ 0 & 1 & 0 & T_y \\ 0 & 0 & 1 & T_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Operation Order

$$\vec{v}' = T \times R \times \vec{v}$$

rotate, then translate

$$\vec{v}' = R \times T \times \vec{v}$$

rotate around a point

$$\vec{v}' = \vec{v} \times T \times R$$

column-major

$$\vec{v}' = \vec{v} \times R \times T$$

$$\vec{v}' = \vec{v} \times T \times R$$

$$\vec{v}' = T \times R \times \vec{v}$$

row-major

Where R is a rotation matrix, T is a translation matrix, v is a vector and v' is the resulting vector.

So, there is a specific order for multiplying matrices together, or vectors by matrices. We also see that 4d matrices are non-commutative - we get a different result if we multiply them in a different order. We will do vector and matrix multiplication in both C and in shaders, so it's worth remembering which convention is being used. In the column-order examples, you can see we are going right-to-left, e.g. input vector, rotate, then translate.

A minor point - in maths notation, row-major vectors are drawn in a row, and column-major vectors are drawn in a column. Matrices are expressed as a capital letter, and vectors a lower-case letter, either in bold-face, or with a harpoon/arrow symbol above them. Normalised (unit) vectors have a "hat" \hat{v} rather than an arrow.

So, what have we learned? We can write a 4X4 matrix directly into a 1d array in C, and send that to a shader using `glUniform`, where it will appear as a uniform mat4. We can then use it to transform our vertex points using an affine transformation. We already have the layout of a translation matrix, so let's try that.

Translation

Let's define a matrix that moves the triangle 0.5 units to the right. Remember that the edge of the view is -1 on the left and 1 on the right, so we don't want to move it completely off the screen. Our matrix will look like this in column-major layout:

```
float matrix[16] {  
    1.0f, 0.0f, 0.0f, 0.0f, // first column  
    0.0f, 1.0f, 0.0f, 0.0f, // second column  
    0.0f, 0.0f, 1.0f, 0.0f, // third column  
    0.5f, 0.0f, 0.0f, 1.0f // fourth column  
};
```

It looks transposed to the way we would write it down in matrix notation when we code it up like this, but I don't want to have to transpose it when I give it to GL.

Now we need to add the uniform variable into the vertex shader, and actually do the multiplication with the vertex point. The vertex shader will look something like this:

```
const char *vertex_shader =  
"#version 430\\n"  
"layout(location = 0) in vec3 vertexPosition;"  
"layout(location = 1) in vec3 vertexColor;"  
  
"uniform mat4 matrix;"  
  
"out vec3 color;"  
  
"void main ()"  
"{"  
"color = vertexColor;"  
"gl_Position = matrix * vec4(vertexPosition, 1.0);"  
"}";
```

Note that the multiplication here is in column-major order as well. Now we should be able to retrieve the location of the new matrix uniform. Do this somewhere after where you link the shader programme, but before entering the main loop. Once we have that, we can "use" the shader programme, and send in our matrix.

We put this code at the end of the load function of Game.cpp :

```
shader.use();  
int matrixLocation = glGetUniformLocation(shader.programId, "matrix");  
glUniformMatrix4fv(matrixLocation, 1, GL_FALSE, matrix);
```

Moving the triangle

Now we will move the triangle back and forth.

Game.h

```
...
private:
...
float speed = 1.0f;
float lastPosition = 0.5f;
float matrix[16] {
    1.0f, 0.0f, 0.0f, 0.0f, // first column
    0.0f, 1.0f, 0.0f, 0.0f, // second column
    0.0f, 0.0f, 1.0f, 0.0f, // third column
    0.5f, 0.0f, 0.0f, 1.0f // fourth column
};
...

```

Game.cpp

```
void Game::update(float dt) {
    if (fabs(lastPosition) > 1.0f) {
        speed = -speed;
    }

    matrix[12] = speed * dt + lastPosition;
    lastPosition = matrix[12];

    shader.use();
    int matrixLocation = glGetUniformLocation(shader.programId, "matrix");
    glUniformMatrix4fv(matrixLocation, 1, GL_FALSE, matrix);
}
```

Rotating the triangle

Your turn ! Use your maths.

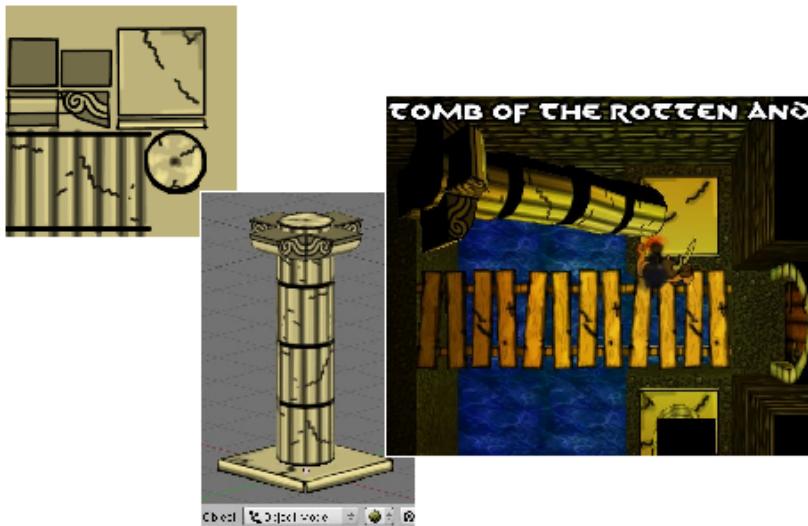
If you want to rotate + scale + translate, you can create your own math library or use GLM : <https://glm.g-truc.net/>

Textures

Texture mapping

Texture mapping is another algorithm from the 1970s that real-time rendering technology wasn't ready for until the 1990s. Edwin Catmull (currently president of Pixar) is credited with creating the idea in 1974. The appropriate reference is his PhD thesis; "Computer display of curved surfaces".

In real-time rendering history, it was adopted in 1990 by the team building Ultima Underworld. There's a nice article about that development on Wikipedia. It was beaten to market by Catacomb 3-D (Id Software) which used their own approach to texturing. The idea is, that instead of using simple colours for surfaces, we will load the surface colours from an image file, and map them to the surface. Not copy, but map - because our surfaces are not going to be the same size as the image, and some maths are required to match them up.



The texturing process; (left) a square image is loaded from a file as a texture. (centre) A designer will assign every vertex on a mesh with a point on the texture; creating a texture coordinate for each vertex. (right)

As the vertices are transformed in 3d, the rendering software will work out the texture coordinate of every fragment on the surface by interpolating the texture coordinates of its surrounding vertices, and sampling the colour of the corresponding point in the image.

We have some special terminology to avoid confusion between on-screen pixels, those from the original image, and the final coloured elements on a surface:

- Texture, an image (usually loaded from a file) that can be mapped onto a surface, and drawn to match the 3d perspective of the surface.
- Pixels, or "picture elements", are the final on-screen, coloured spots rendered to the viewport. Your GL viewport might have 1024x768 pixels.
- Fragments are pixel-sized areas of a surface (the visible surface is divided into pixel-sized 2d fragments). As a 3d triangle gets closer to the camera, its surface will contain more and more fragments. Fragments can overlap, and we can blend them together for techniques like partial-transparency, but there is only 1 final pixel rendered.
- Texels, or "texture elements", are pixels loaded from an image. If we map a 512x512 texture to a surface that occupies only 30 pixels of the viewport, we have more texels than fragments.

To load textures, we will use `SDL2_Image`.

Loading texture

Create a class for the texture.

src/engine/Texture.h

```
#ifndef TEXTURE_H
#define TEXTURE_H

#include <SDL.h>
#include <string>
using std::string;

class Texture
{
public:
    Texture();
    ~Texture();

    void unload();
    bool load(const string& filenameP);
    void updateInfo(int& widthOut, int& heightOut);
    void use() const;

    unsigned int getTextureID() const { return textureId; }
    inline int getWidth() const { return width; }
    inline int getHeight() const { return height; }

private:
    unsigned int textureId;
    string filename;
    int width;
    int height;
};

#endif
```

src/engine/Texture.cpp

```
#include "Texture.h"

#include <GL/glew.h>
#include <SDL_image.h>

#include <iostream>
#include "Log.h"

Texture::Texture() : textureId(0), filename(""), width(0), height(0) {}

Texture::~Texture() {}

void Texture::unload() {
    glDeleteTextures(1, &textureId);
}

bool Texture::load(const string& filenameP) {
    filename = filenameP;
    // Load from file
    SDL_Surface* surf = IMG_Load(filename.c_str());
    if (!surf) {
        LOG(Error) << "Failed to load texture file " << filename;
        return false;
    }
    width = surf->w;
    height = surf->h;
    int format = 0;
    if (surf->format->format == SDL_PIXELFORMAT_RGB24) {
        format = GL_RGB;
    } else if (surf->format->format == SDL_PIXELFORMAT_RGBA32) {
        format = GL_RGBA;
    }

    // Generate texture
    glGenTextures(1, &textureId);
    glBindTexture(GL_TEXTURE_2D, textureId);
    glTexImage2D(GL_TEXTURE_2D, 0, format, width, height, 0, format, GL_UNSIGNED_BYTE, surf->pixels);
    SDL_FreeSurface(surf); // We no longer need the surface

    // Generate mipmaps for texture
    glGenerateMipmap(GL_TEXTURE_2D);
    // configure wrapping
    glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_CLAMP_TO_EDGE);
    glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP_TO_EDGE);
    // Enable linear filtering
    glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR_MIPMAP_LINEAR);
    glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR);

    LOG(Info) << "Loaded texture " << filename;
    return true;
}

void Texture::updateInfo(int& widthOut, int& heightOut) {
    widthOut = width;
    heightOut = height;
}

void Texture::use() const { glBindTexture(GL_TEXTURE_2D, textureId); }
```

Copy image data to OpenGL texture

Create the texture

We need to generate a new texture object, which will give us an unsigned integer to refer to it with later. Before we can copy data into the texture we have to bind it into focus. Any time that we bind a texture it will be moved into one of several active texture slots. We don't need to worry about this yet, but it's good to specify a slot now, otherwise it will use the most recently activated slot, which might create unexpected interference with another part of your programme - remember GL is a state machine. . This is done in the load function :

```
...
// Generate texture
 glGenTextures(1, &textureID);
 glBindTexture(GL_TEXTURE_2D, textureID);
 glTexImage2D(GL_TEXTURE_2D, 0, format, width, height, 0, format, GL_UNSIGNED_BYTE, surf->pixels);
 SDL_FreeSurface(surf); // We no longer need the surface

// Generate mipmaps for texture
 glGenerateMipmap(GL_TEXTURE_2D);
// configure wrapping
 glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_CLAMP_TO_EDGE);
 glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP_TO_EDGE);
// Enable linear filtering
 glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR_MIPMAP_LINEAR);
 glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR);
...

```

Look up the `glTexImage2D` function in the official documentation. It has a lot of parameters, but all of them should relate to how we loaded our image. At this stage we are just using bog-standard 2d textures, but we have to specify this whenever we bind or create a texture anyway. We aren't going to set up any custom levels-of-detail texture, so we just put 0 for this parameter. There are lots of options for texture formats - it's a good idea to have a look at the tables of options. We know for sure that our loaded image is RGB or RGBA format now, so we can say `GL_RGB` or `GL_RGBA` in the `format` variable. We give it the width and height in pixels, which were returned by our image loader. The next "0" is an unused parameter. Again, we can put `format` as our internal format. It's possible to have the bytes of each pixel ordered differently. Finally, we know that our image data is in `unsigned char` (byte) format, and we give it the pointer to our data. The last four lines before the return are a good, safe, default "wrapping" mode to use for loaded textures, and a good default for anti-aliasing. We'll get onto both of these topics later.

Active Texture Slots

The texture is now loaded, and we can use it whenever we want to render by activating a texture slot, and binding the texture. It will stay in that slot until another texture is bound into it, so if your programme only has 1 texture then you can just do that once.

OpenGL can load a large number of textures into memory, but your GPU can only read from a small number of textures at one time (most GPUs will read 8 or so). To manage this, the state machine has a number of what it calls "active textures". We need to swap our textures in and out of these active texture slots as we render different objects. Yes, we have to do this low-level tedium ourselves, and it's very easy to make mistakes. The default active texture slot is number 0.

Texture on the shader

UV/texture coordinates

We will pass to our vertex shader the coordinates we need to draw the texture on two triangles. They will be forwarded to the fragment shader :

```
const char* vertexShader =
"#version 430\n"
"layout (location = 0) in vec3 vertexPosition;"
"layout (location = 1) in vec2 uv;"

"out vec2 textureCoordinates;"

"void main ()"
"{"
    "textureCoordinates = uv;"
    "gl_Position = vec4(vertexPosition, 1.0);"
}"

const char* fragmentShader =
"#version 430\n"

"layout (binding = 0) uniform sampler2D basicTexture;"

"in vec2 textureCoordinates;"
"out vec4 fragColor;"

"void main ()"
"{"
    "vec4 texel = texture(basicTexture, textureCoordinates);"
    "fragColor = texel;"
}"
```

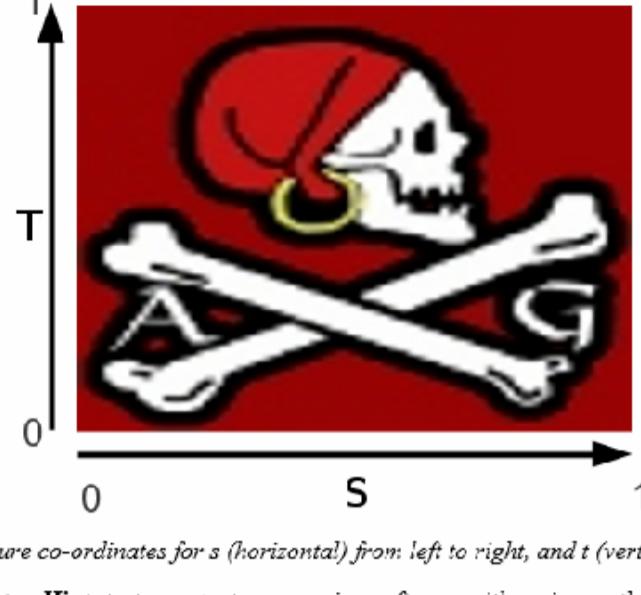
The fragment shader will use those texture coordinates and a uniform to get the data from the texture file. The type of this data is sampler2D.

Once we have those two data, we can use the GLSL texture() function to deduce the fragment color.

Texture in fragment shader

Texture coordinates

The main problem is that the surfaces are not going to have the same number of fragments as the image has texels, and this will change as the surface is viewed at viewpoints and perspectives. The obvious solution is to calculate a very simple interpolation function for each fragment of the wall, to find out the closest-matching texel from the image.



GL texture co-ordinates for s (horizontal) from left to right, and t (vertical) from bottom to top. Hint: test your texture mapping software with an image that has text in it; this will help you spot if you've got one of the axes reversed.

To keep interpolation of texels independent of the image size, we use texture coordinates, which are a horizontal and vertical value between 0.0 and 1.0. In OpenGL 0.0 is the left or bottom of a texture, and 1.0 is the right, or top of a texture. Direct3D has 0.0 at the top, and 1.0 at the bottom - don't mix them up!

So we can make another per-vertex variable to keep track of what part of the image each vertex should map to; a texture coordinate. This will be 2 floats for each vertex, and we will make another vertex buffer to store these. Our previous vertex buffers have all had 3 floats - don't accidentally tell GL that it is 3d, or it will get the memory "stride" (ordering) wrong, and you'll have mixed-up values.

We use x,y,z,w to refer to 3d space, and r,g,b,a for colours. In most 3d libraries u,v are used to refer to 2d space, but OpenGL prefers s,t. The GLSL language vector data-types (vec2, vec3, vec4) have built-in short-cuts called swizzle operators. You may have used something like vec3 pos = result.xyz; or float red = Kdr; already. You can also use .s, .t, and .st.

Creating Texture Coordinates

For a wall mesh, made up of 2 triangles, we will create 18 vertices and 6 texture coordinates, and put them in a vertex buffer.

Game.cpp

```
...
void Game::load() {
    GLfloat points[] = {
        0.5f, 0.5f, 0.0f,
        0.5f, -0.5f, 0.0f,
        0.5f, 0.5f, 0.0f,
        0.5f, 0.5f, 0.0f,
        -0.5f, 0.5f, 0.0f,
        -0.5f, -0.5f, 0.0f,
        0.5f, 0.5f, 0.0f,
        0.5f, 0.5f, 0.0f,
        -0.5f, 0.5f, 0.0f,
        -0.5f, -0.5f, 0.0f
    };

    GLfloat texcoords[] = {
        0.0f, 1.0f,
        0.0f, 0.0f,
        1.0f, 0.0f,
        1.0f, 0.0f,
        0.0f, 1.0f,
        0.0f, 0.0f,
        1.0f, 0.0f,
        1.0f, 0.0f,
        0.0f, 1.0f,
        0.0f, 0.0f,
        1.0f, 0.0f,
        1.0f, 0.0f
    };

    GLuint pointsVbo = 0;
    glGenBuffers(1, &pointsVbo);
    glBindBuffer(GL_ARRAY_BUFFER, pointsVbo);
    glBufferData(GL_ARRAY_BUFFER, sizeof(points), points, GL_STATIC_DRAW);

    GLuint texVbo;
    glGenBuffers(1, &texVbo);
    glBindBuffer(GL_ARRAY_BUFFER, texVbo);
    glBufferData(GL_ARRAY_BUFFER, sizeof(texcoords), texcoords, GL_STATIC_DRAW);

    GLuint vao = 0;
    glGenVertexArrays(1, &vao);

    glBindVertexArray(vao);
    glBindBuffer(GL_ARRAY_BUFFER, pointsVbo);
    glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 0, NULL);

    glBindBuffer(GL_ARRAY_BUFFER, texVbo);
    int dimensions = 2; // 2d data for texture coords
    glVertexAttribPointer(1, dimensions, GL_FLOAT, GL_FALSE, 0, NULL);

    glEnableVertexAttribArray(0);
    glEnableVertexAttribArray(1);

    Texture texture;
    texture.load("assets/character.png");
    texture.use();

    shader.compileVertexShader();
    shader.compileFragmentShader();
    shader.createShaderProgram();

    shader.use();
    int texLoc = glGetUniformLocation(shader.programId, "basicTexture");
    glUniform1i(texLoc, 0); // use active texture 0
}
```

It would be optimum to mix vertices and texture coordinates, and to use an other draw mode, we will do that later.

Don't forget to draw the two triangles after biding the VAO :

```
glDrawArrays(GL_TRIANGLES, 0, 6);
```

You also need to change the direction of the culling in Window.cpp, or your texture won't appear :

```
...
// Enable cull face optimization
glEnable(GL_CULL_FACE); // Cull face
glCullFace(GL_BACK); // Cull back face
glFrontFace(GL_CCW); // GL_CCW for Counter Clock-Wise
...
```

Example : A Breakout Game

Use your OpenGL framework and OOP to implement the example breakout game :
<https://learnopengl.com/In-Practice/2D-Game/Breakout>